S3 (Simple Storage Service)

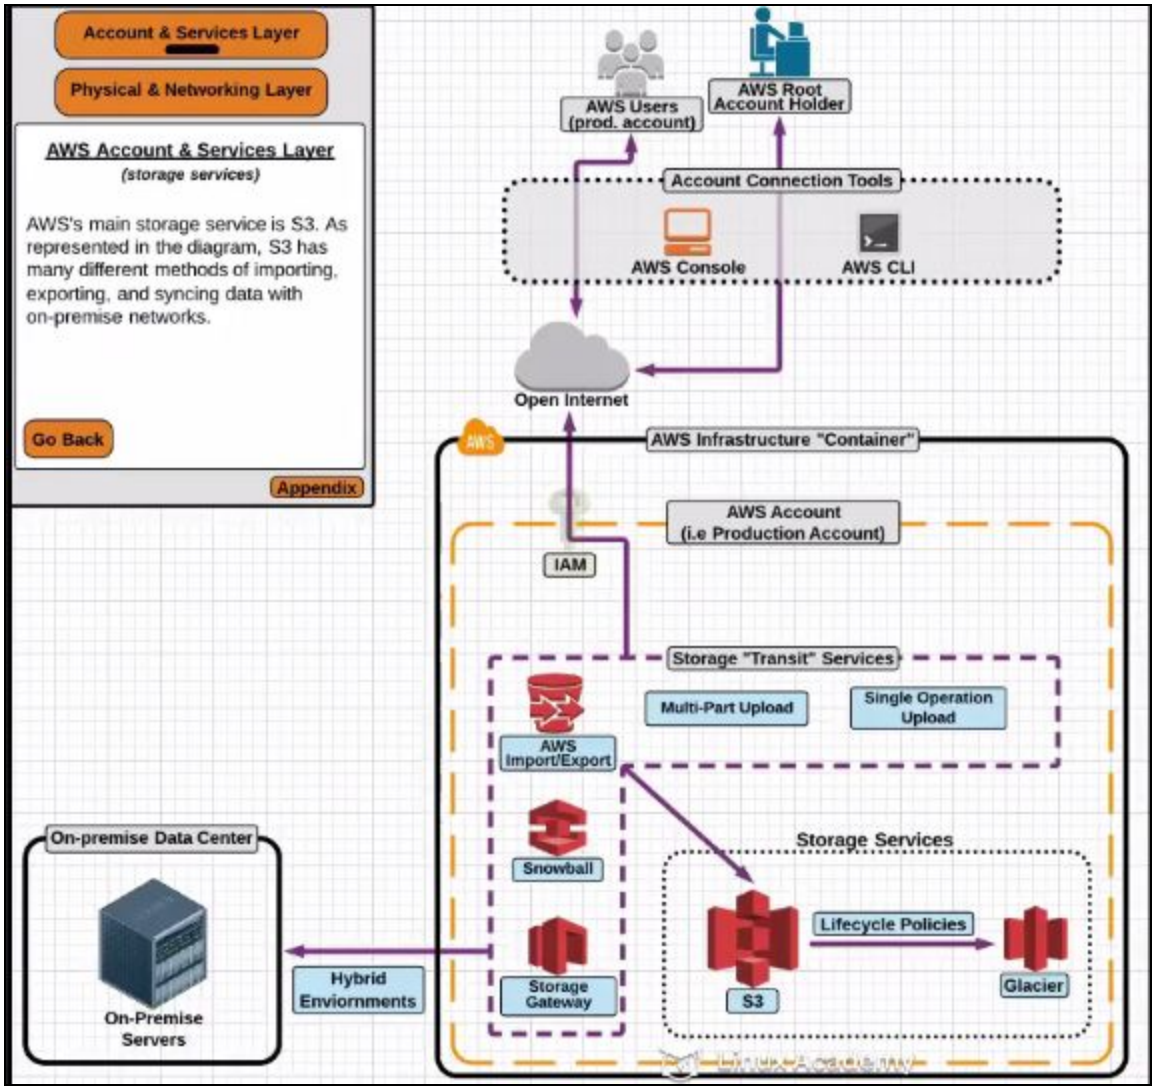

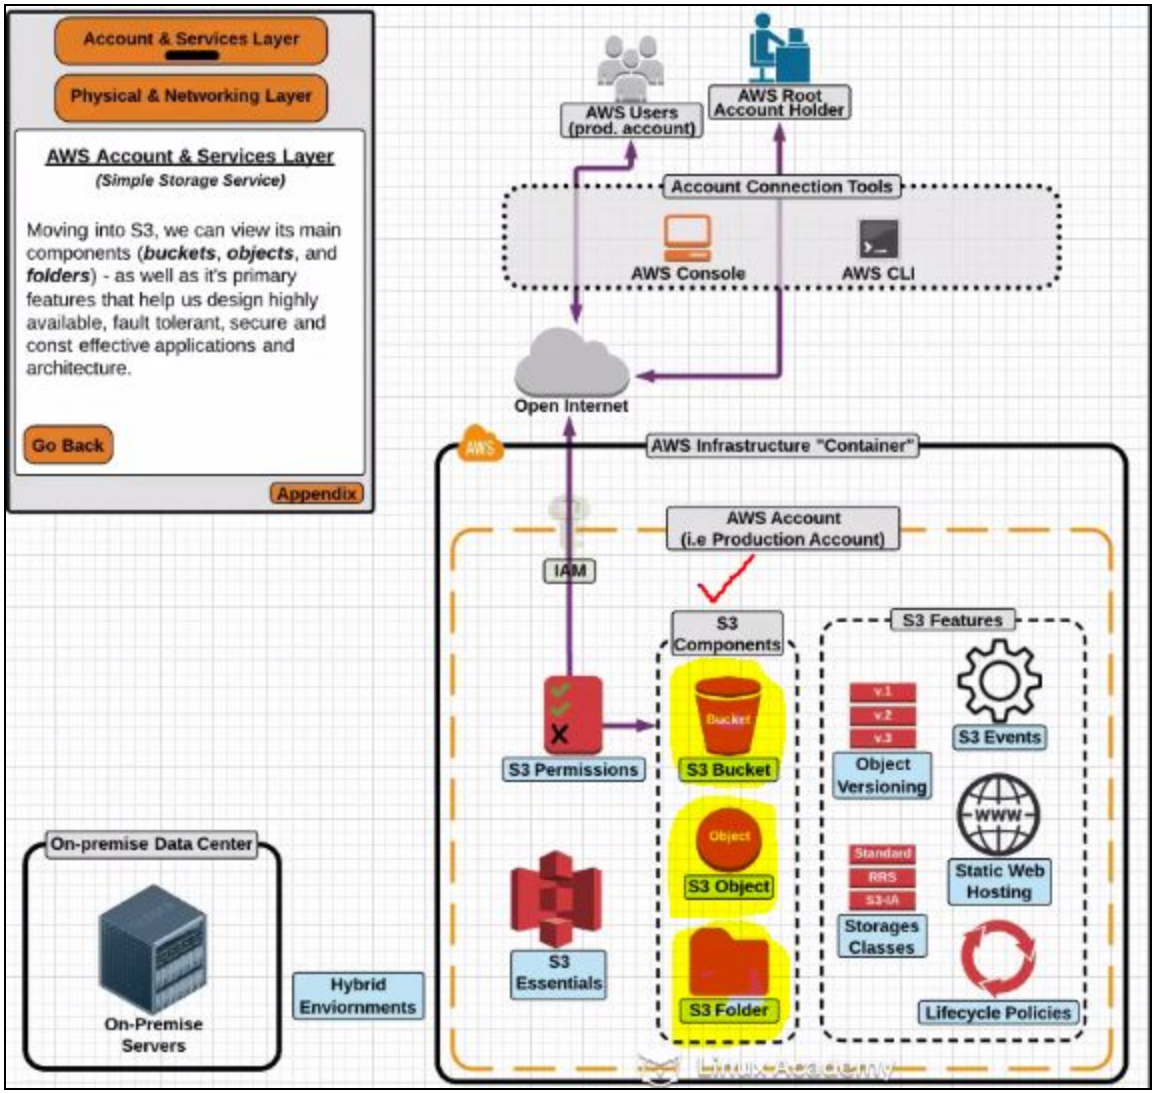
Account & Services Layer
Physical & Networking Layer
AWS Account & Services Layer
(Simple Storage Service)
Moving into S3, we can view its main components (buckets, objects, and folders) - as well as it's primary features that help us design highly available, fault tolerant, secure and const effective applications and architecture.
Go Back
Appendix
AWS Users (prod. account)
AWS Root Account Holder
Account Connection Tools
AWS Console
AWS CLI
Open Internet
AWS Infrastructure "Container"
AWS Account (i.e Production Account)
IAM
S3 Components
S3 Features
S3 Permissions
S3 Bucket
S3 Object
S3 Folder
S3 Essentials
Object Versioning
S3 Events
Static Web Hosting
Storages Classes
Lifecycle Policies
On-premise Data Center
On-Premise Servers
Hybrid Enviornments

S3 Essentials

S3 basics (Simple Storage Service)

- Amazon S3 is storage for the internet.
- Amazon S3 is a object based storage not a block-based storage. It means you cannot install apps in this storage, only you can keep files.
- Amazon S3 provides a simple web service interface that you can use to store and retrieve any amount of data, at any time, from anywhere on the internet.

Eg: Using this web service, developers can easily build applications that make use of internet storage. Since S3 is highly scalable and you can pay for what you use, developers can start with small and grow their application as they wish, with no compromise on performance or reliability

- AWS S3 is the main storage service which serves many purposes when designing ***highly available, fault tolerant, highly scalable, reliable, low-latency*** data storage infrastructure at ***very low-cost***.
- S3 is also known as key-value data store (key,value, versionID, MetaData , ACL and sub-resources)

Note: Example of S3 URL : https://s3-ap-south1.amazonaws.com/mybucket1
If upload is successful, then status of returned code is HTTP200

Use Cases:

- Bulk (basically unlimited) static object storage
- Hosting static files & websites
- Object versioning
- Origin for CloudFront CDN
- And many more

Important S3 Facts

- Objects stay within an AWS region and are synced across all AZ's for extremely High Availability and Durability
- You should always create an S3 bucket in a region that makes sense to its purpose:
    - service content to customer
    - sharing data with EC2

Benefits of S3

- 99.99% availability
- 99.999999999% durability (called as eleven 9's durability)

S3 Read Consistency Rule

- ALL regions now support read-after-write consistency for PUTS of new objects into S3
  i.e it means objects are immediately available after putting it in S3 bucket
- All regions use eventual consistency for PUTS overwriting existing objects and DELETES of object
  i.e If you delete or overwrite then there may be some delay

S3 Bucket

S3 Buckets

- Buckets are the main storage container of S3, and contain a grouping of information and have sub namespaces that are similar to folders (but yet are called folders)
- Tags can be used to organize buckets (i.e tags based on application the buckets belongs to)
- **Each bucket must have a unique name across ALL of AWS. Also it should be in small case letters.**

S3 bucket limitations

- Only 100 buckets can be created
- Bucket ownership cannot be changed later also

NOTE: S3 as a service is actually global, you can view all the buckets that are in different regions at the same place (global)

LAB 1:Creating a Bucket

Login to management console and search for S3 service

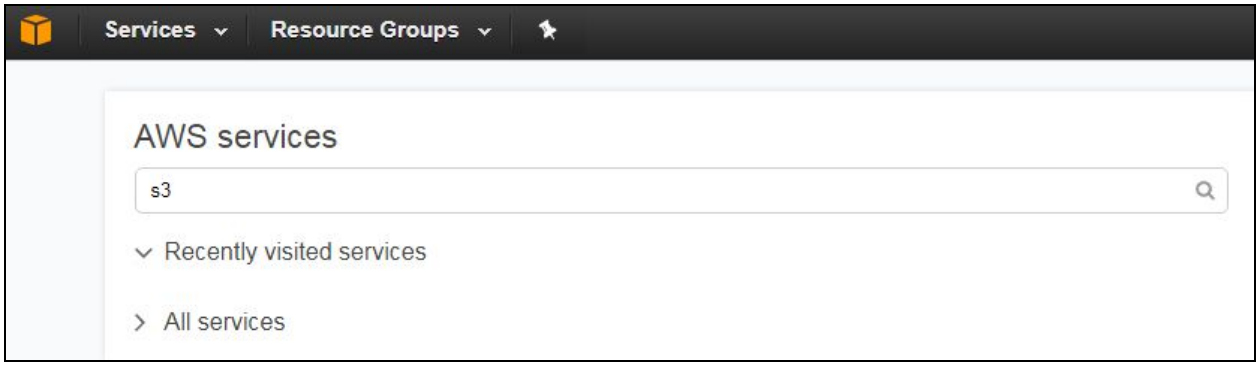

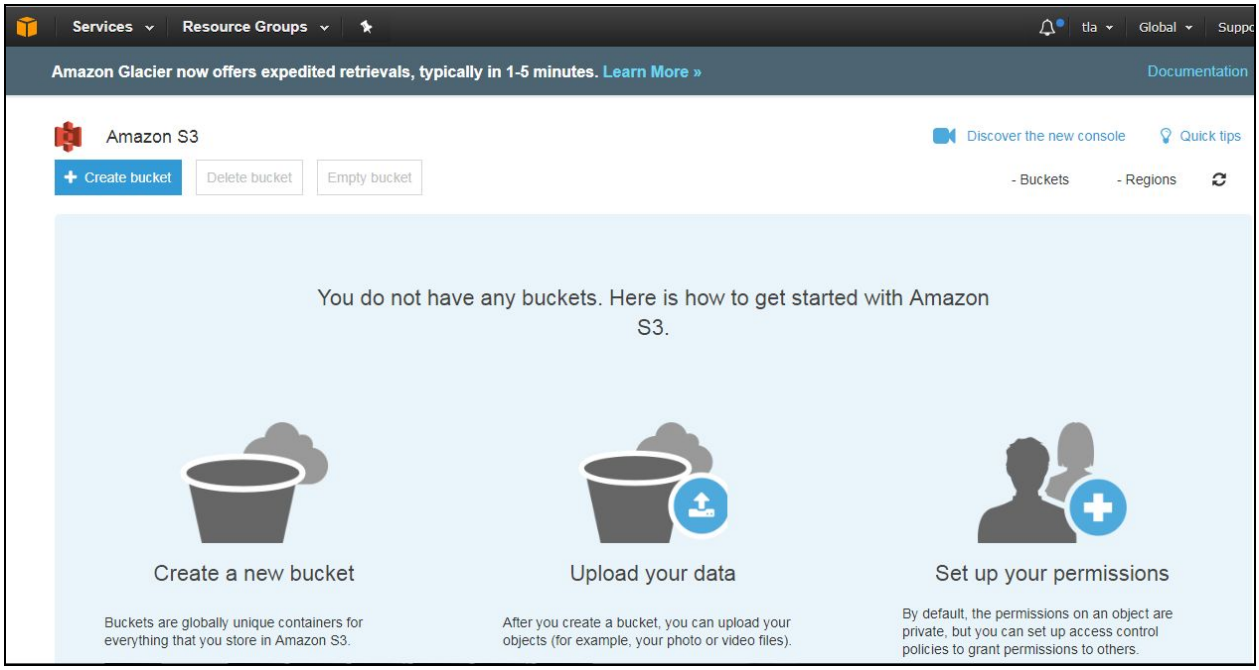

Click on **Create bucket**

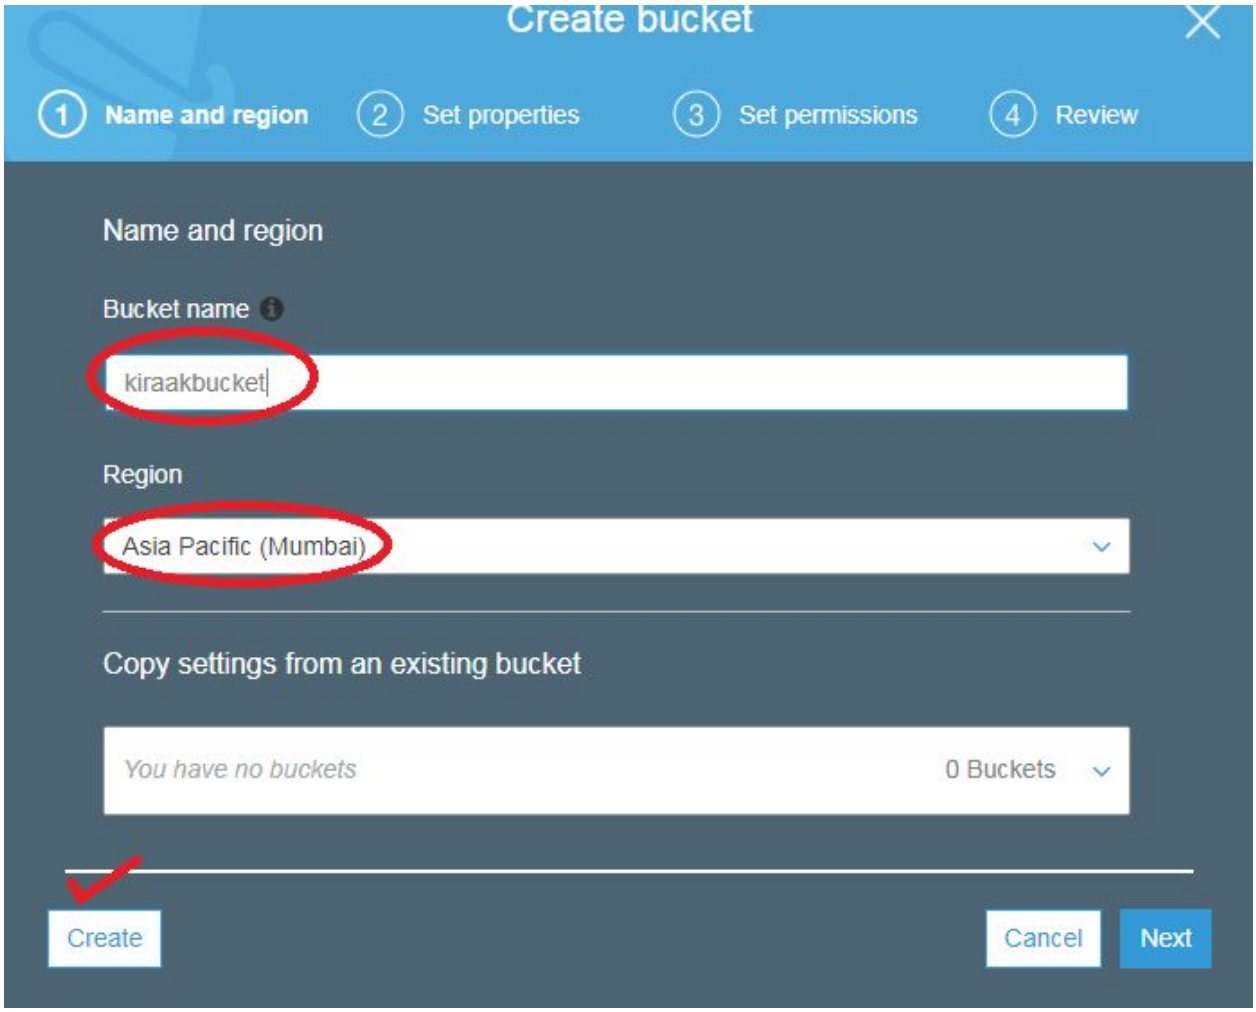

Fill out the below fields
Bucket name:
Region:
Then Click on **Create**

Or you can do the below all 4 tabs

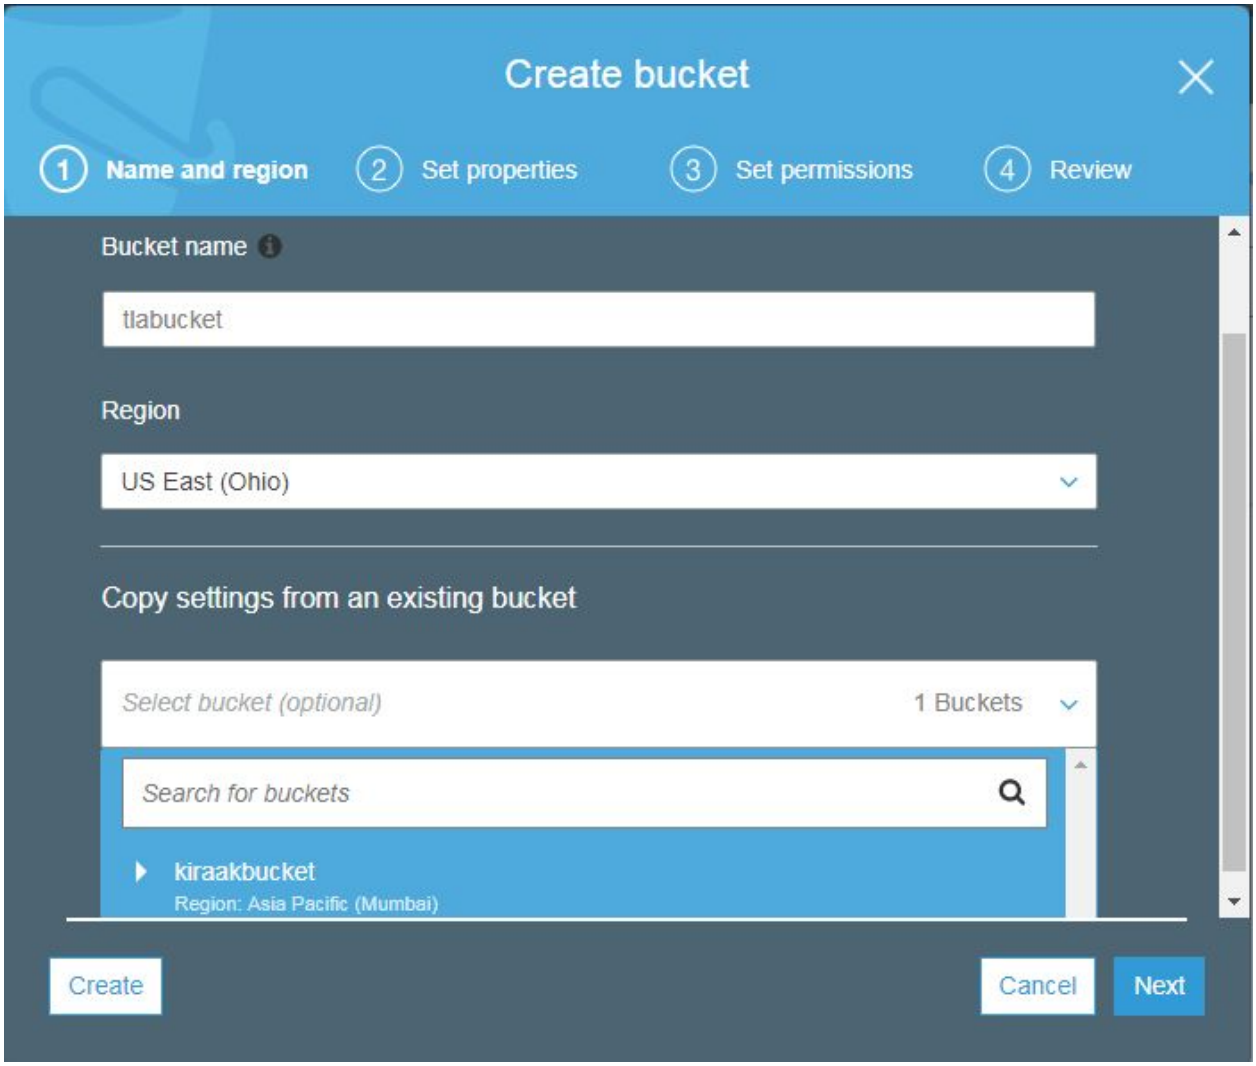

Click on Next , then you will see **properties**

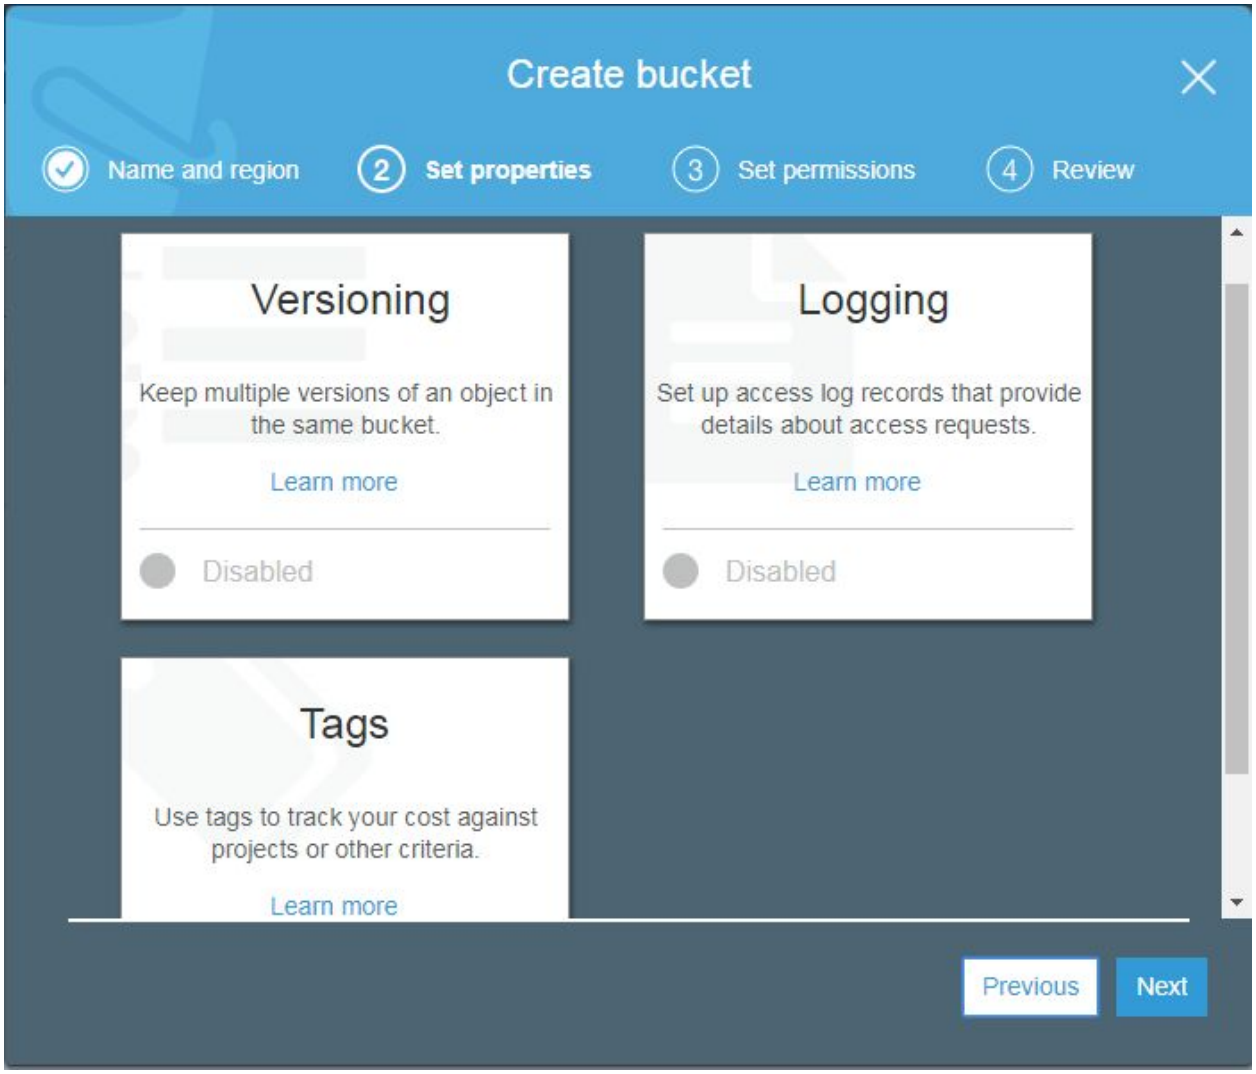

Click on Next, then you will see **Permissions**

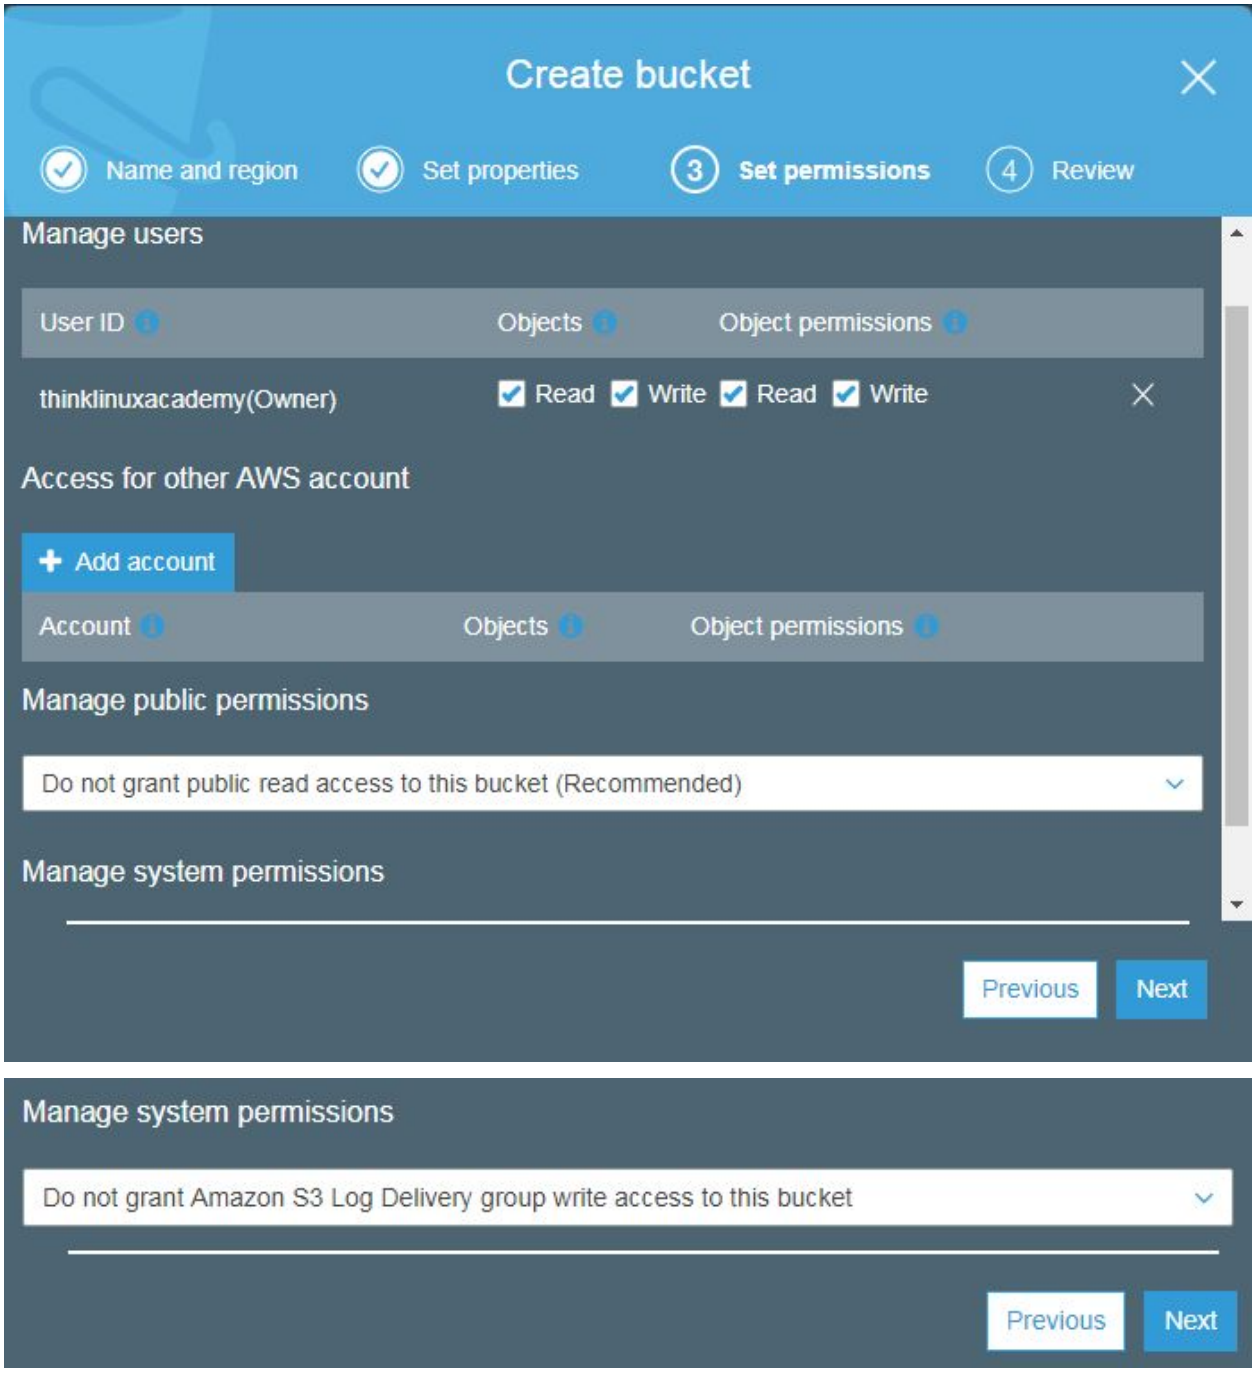

Click on Next, then you will see **Review screen**

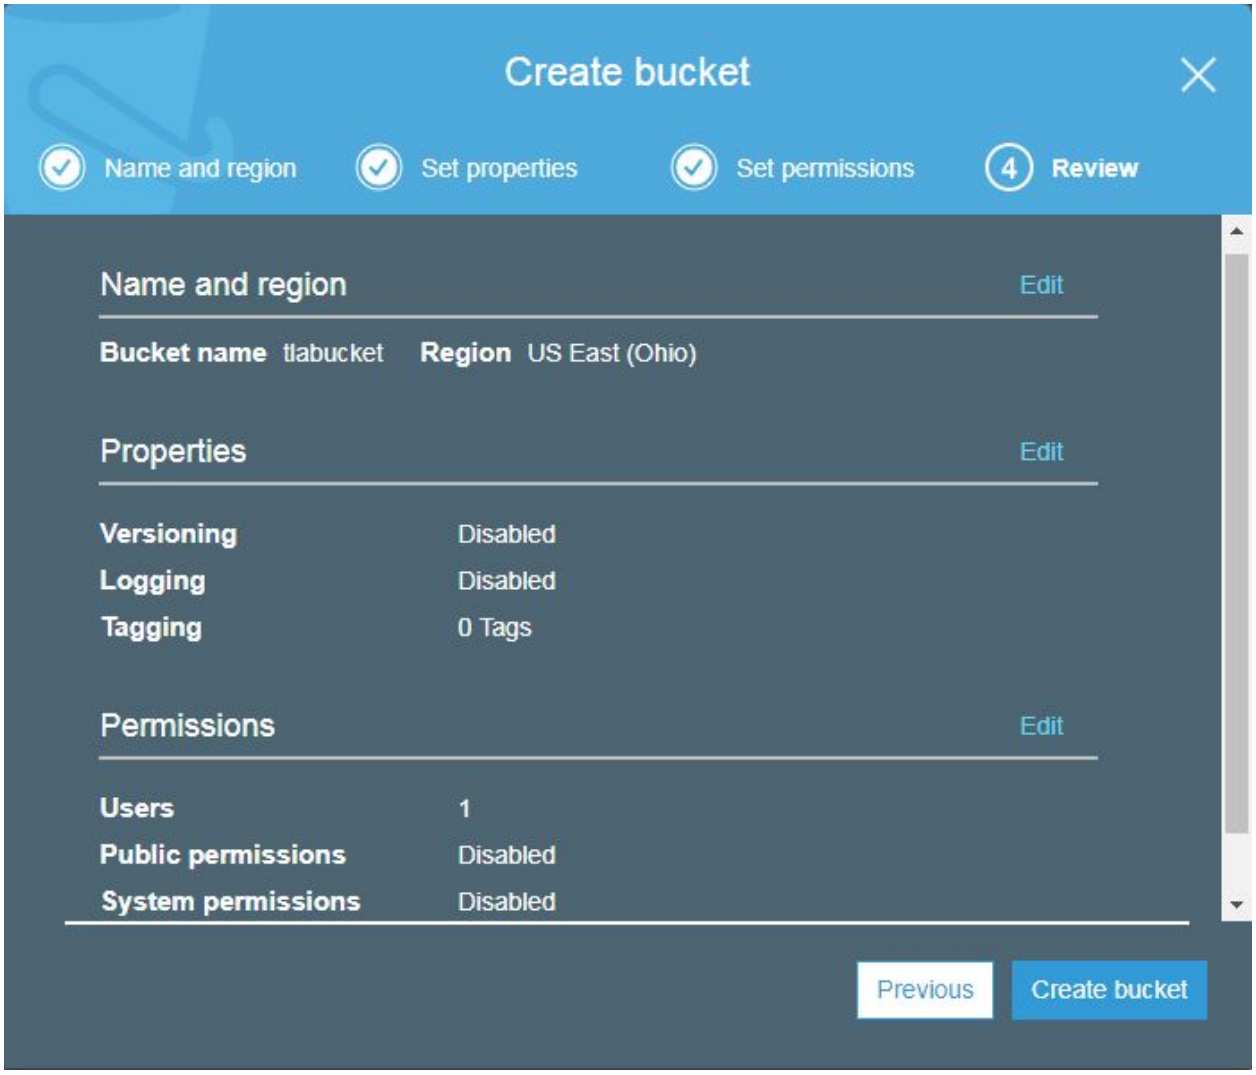

Then click on Create bucket, THE BUCKET IS CREATED

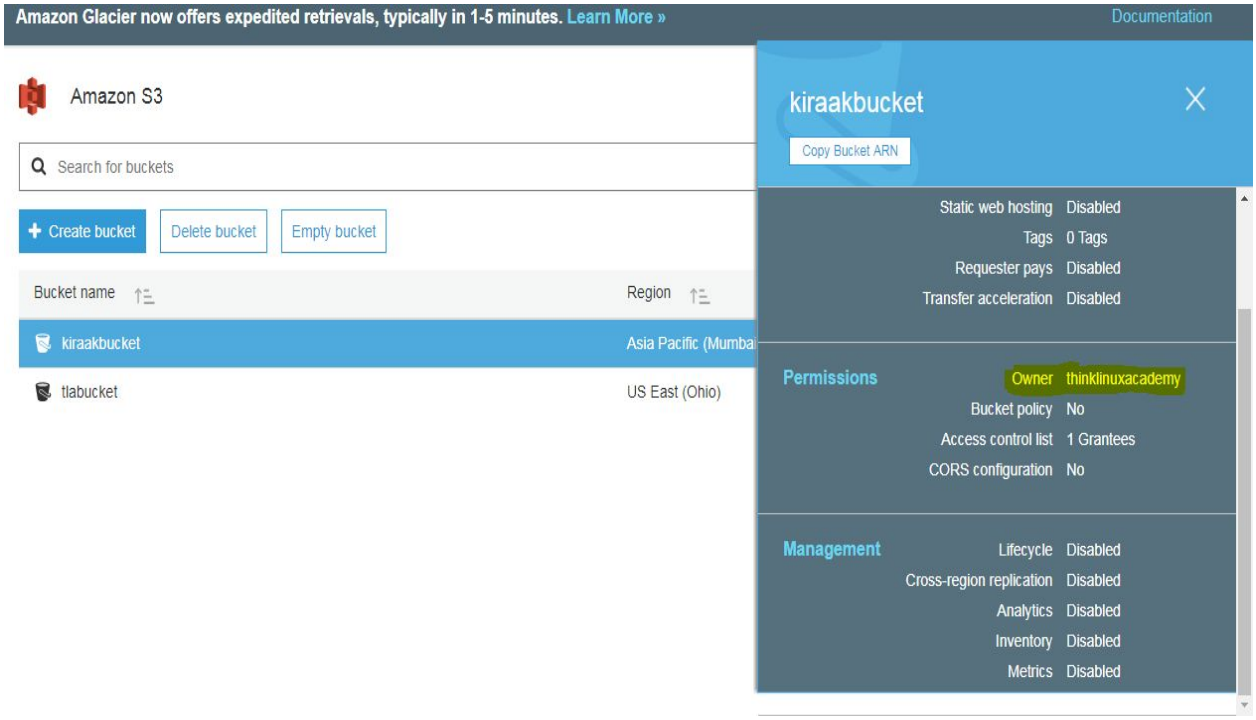

You can see the ownership of the bucket

S3 Objects

S3 Objects

- Objects are static files that contain metadata information :
- Set of name-key pairs
- Contain information specified by the user, and AWS information such as storage type
- Each object must be assigned a **storage type,** which determines the object's **availability, durability and cost**
- By default, all objects are private (means not public)

Object facts

- Be as small as 0 bytes and as large as 5TB
- Can have multiple versions (if versioning is enabled)
- Be made publicly available via a URL
- Automatically switch to a different storage class or deleted (via lifecycle policies)
- Encrypted
- Organized into "sub-name" spaces called folders

Object Encryption

- SSE (Server Side Encryption)
- S3 can encrypt the object before saving it on the partitions in the data centers and decrypt it when it is downloaded
- AES-256
- SSE-KMS
- Also you can use your own encryption keys
- Considered client side encryption where you encrypt the data before upload

- SSL terminated endpoints for the API

LAB 2 : Create a object in the Bucket

First go into the bucket → then upload object (suppose one photo)

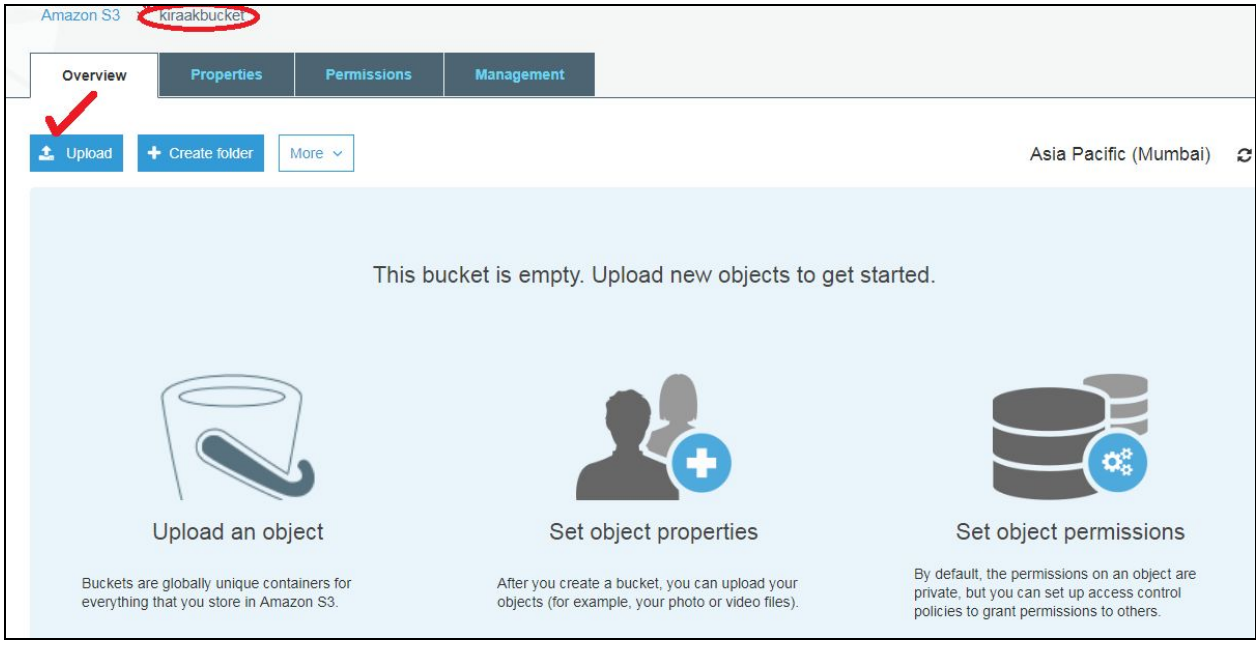

Click on **Upload,** then browse and select the photo

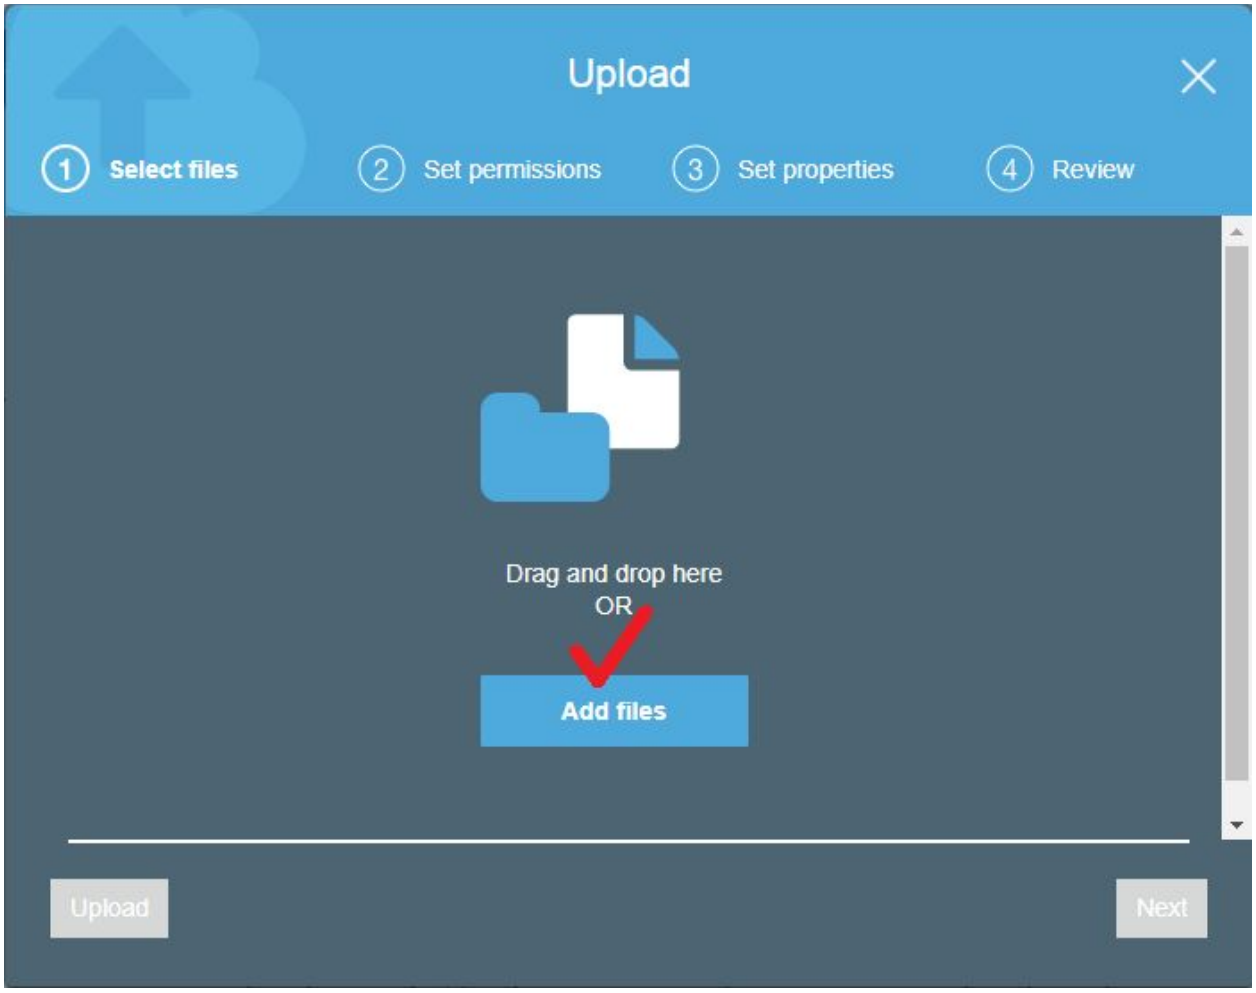
Upload
1 Select files
2 Set permissions
3 Set properties
4 Review
Drag and drop here
OR
Add files
Upload
Next

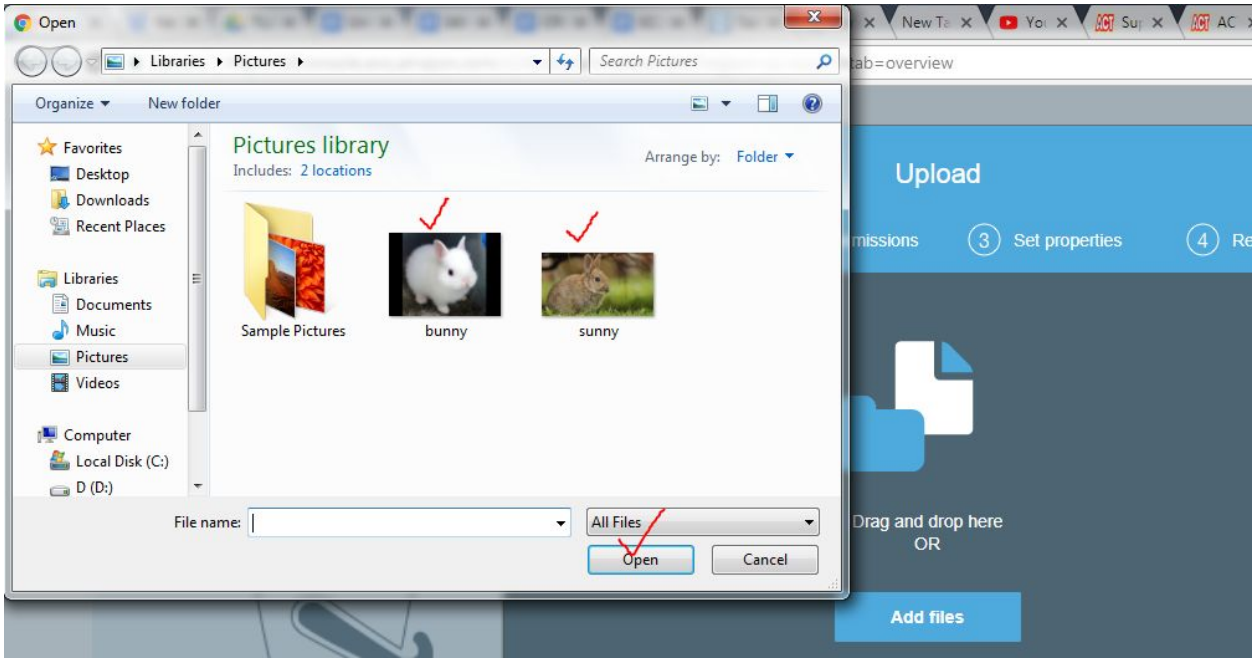
Open
Libraries ▸ Pictures ▸
Search Pictures
Organize
New folder
Favorites
Desktop
Downloads
Recent Places
Libraries
Documents
Music
Pictures
Videos
Computer
Local Disk (C:)
D (D:)
Pictures library
Includes: 2 locations
Arrange by: Folder
Sample Pictures
bunny
sunny
File name:
All Files
Open
Cancel
Upload
Set properties
Drag and drop here
OR
Add files

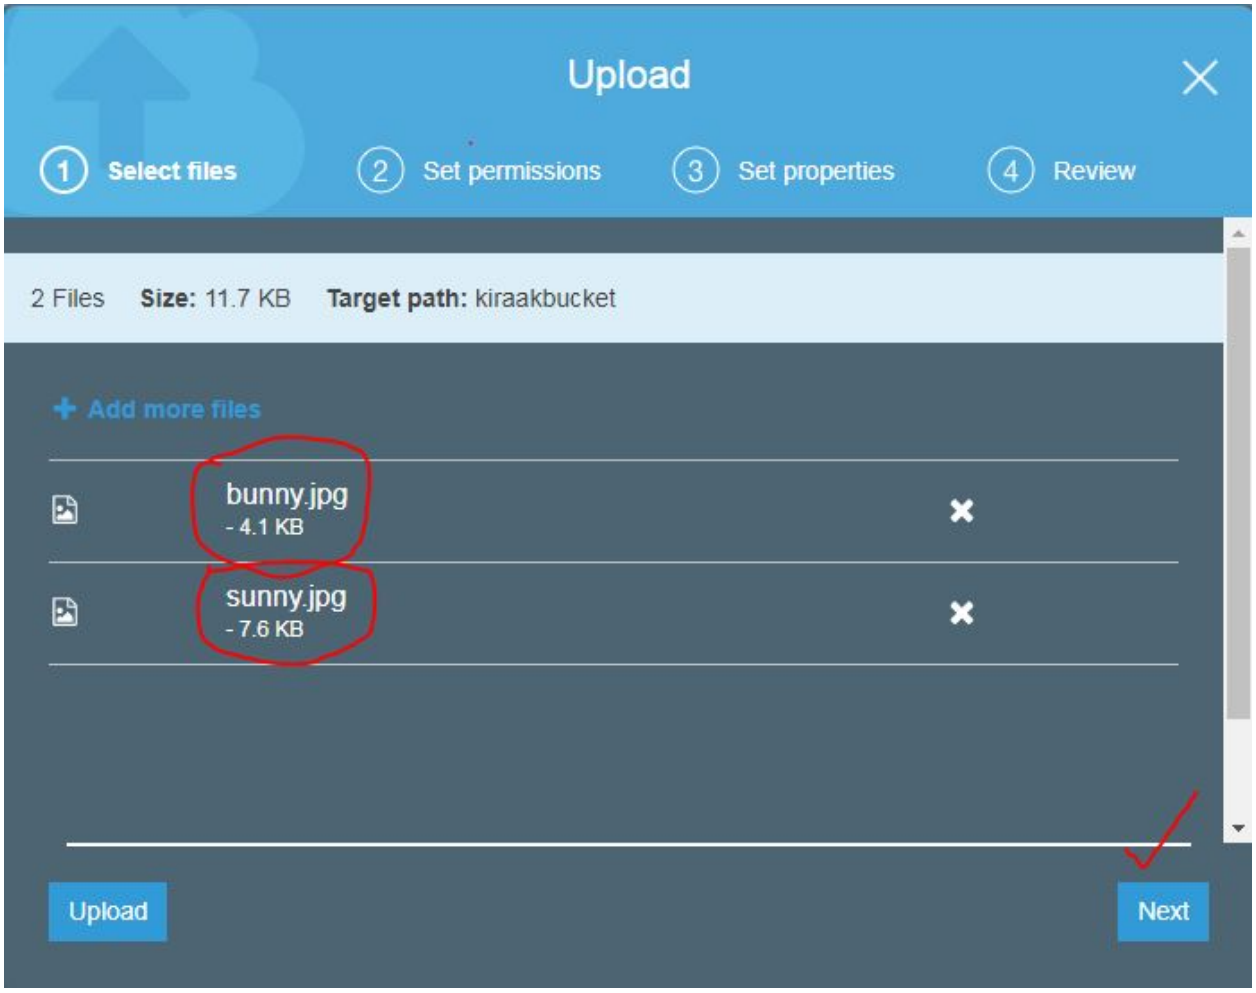
Upload
1 Select files
2 Set permissions
3 Set properties
4 Review
2 Files Size: 11.7 KB Target path: kiraakbucket
+ Add more files
bunny.jpg
- 4.1 KB
sunny.jpg
- 7.6 KB
Upload
Next

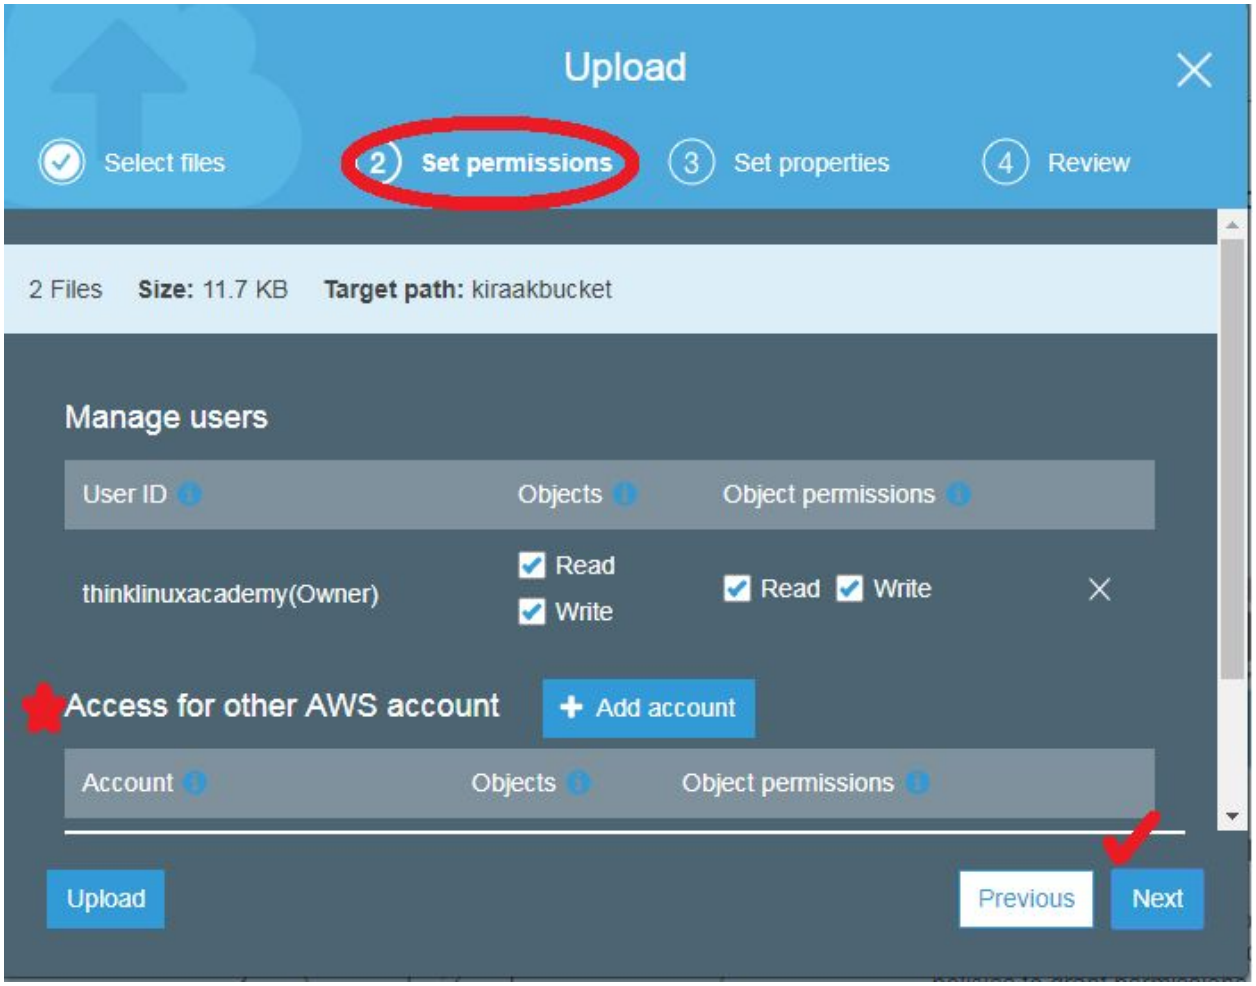
Upload
Select files
2 Set permissions
3 Set properties
4 Review
2 Files Size: 11.7 KB Target path: kiraakbucket
Manage users
User ID
Objects
Object permissions
thinklinuxacademy(Owner)
Read
Write
Read Write
Access for other AWS account
Add account
Account
Objects
Object permissions
Upload
Previous
Next

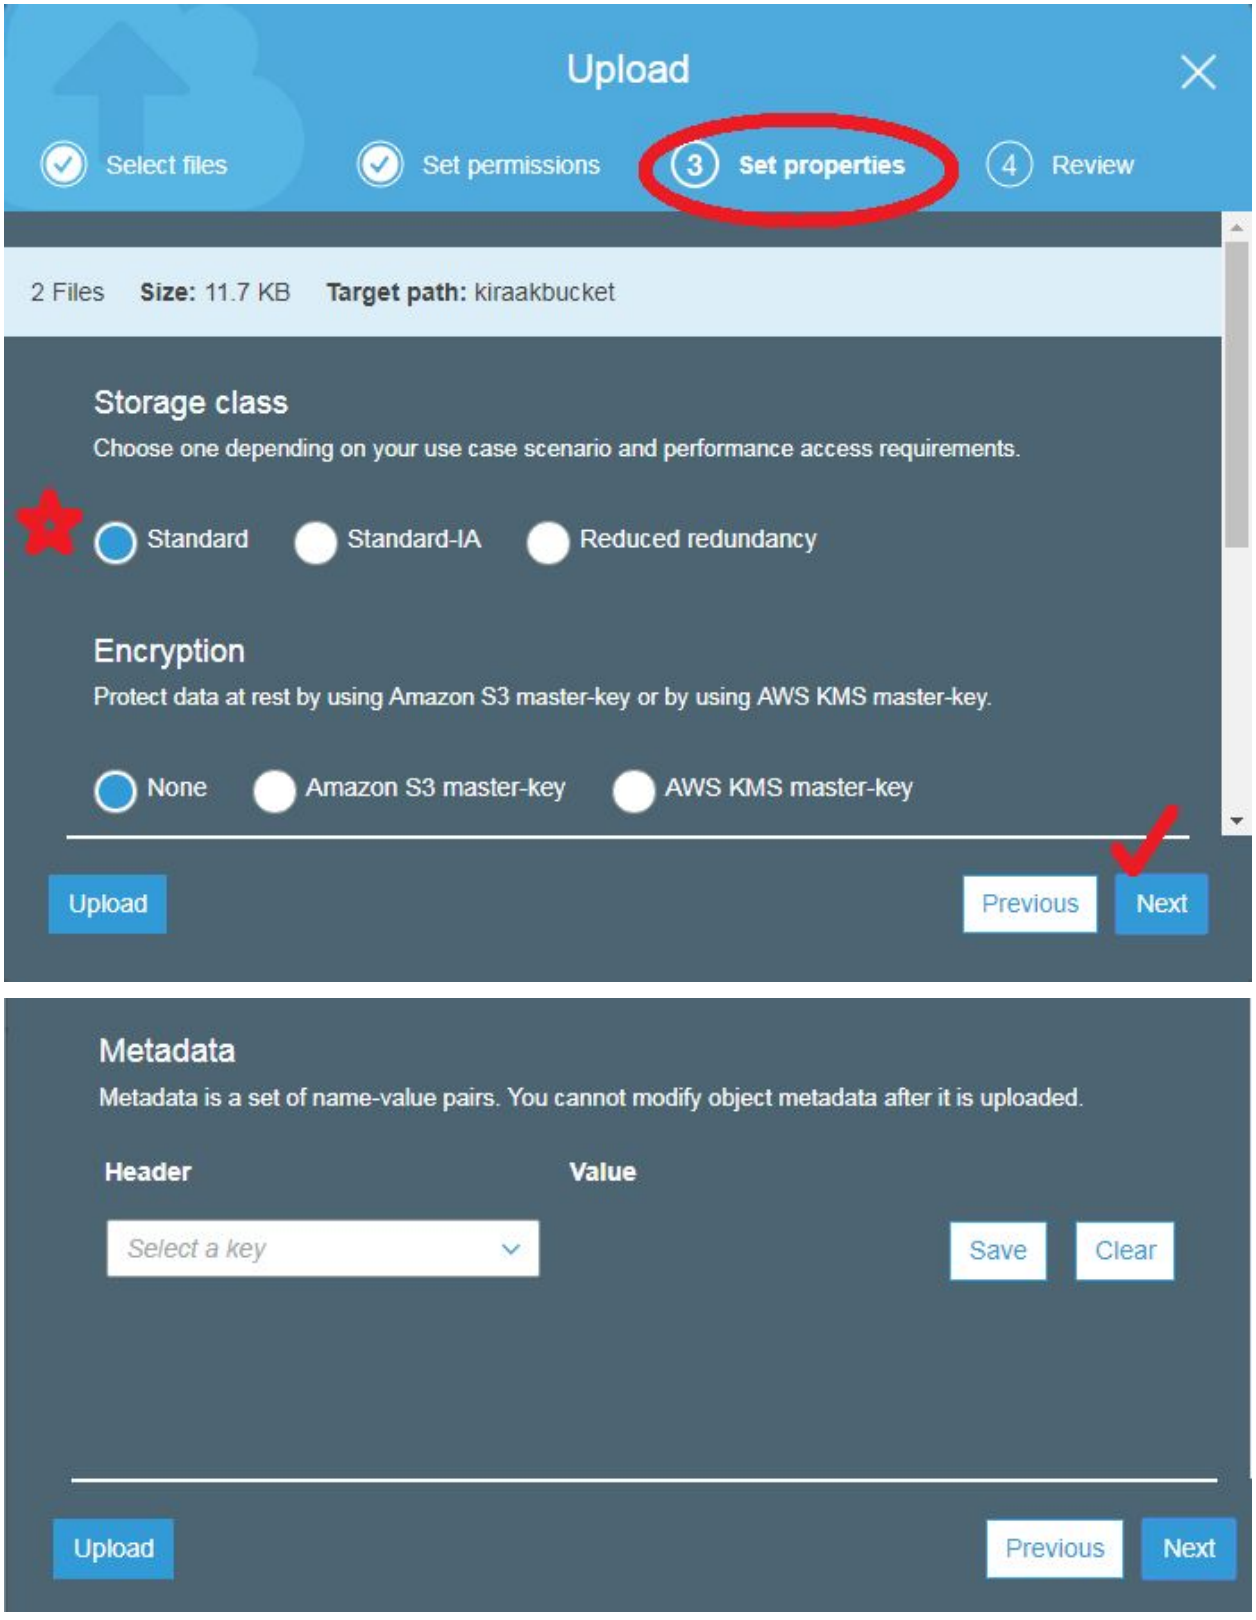

Don’t select metadata, click on Next

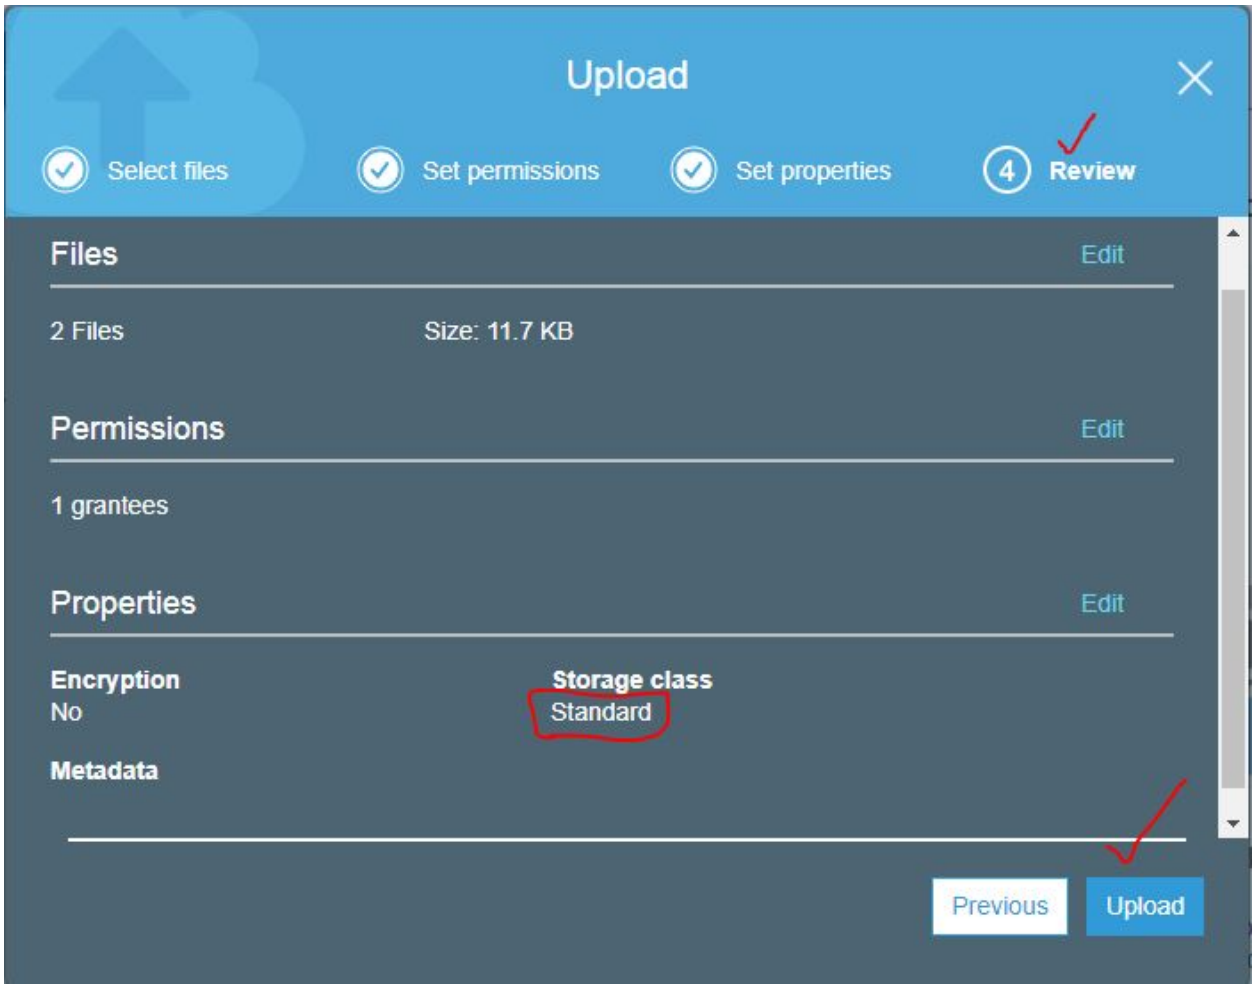

So now you can see the bucket and its content

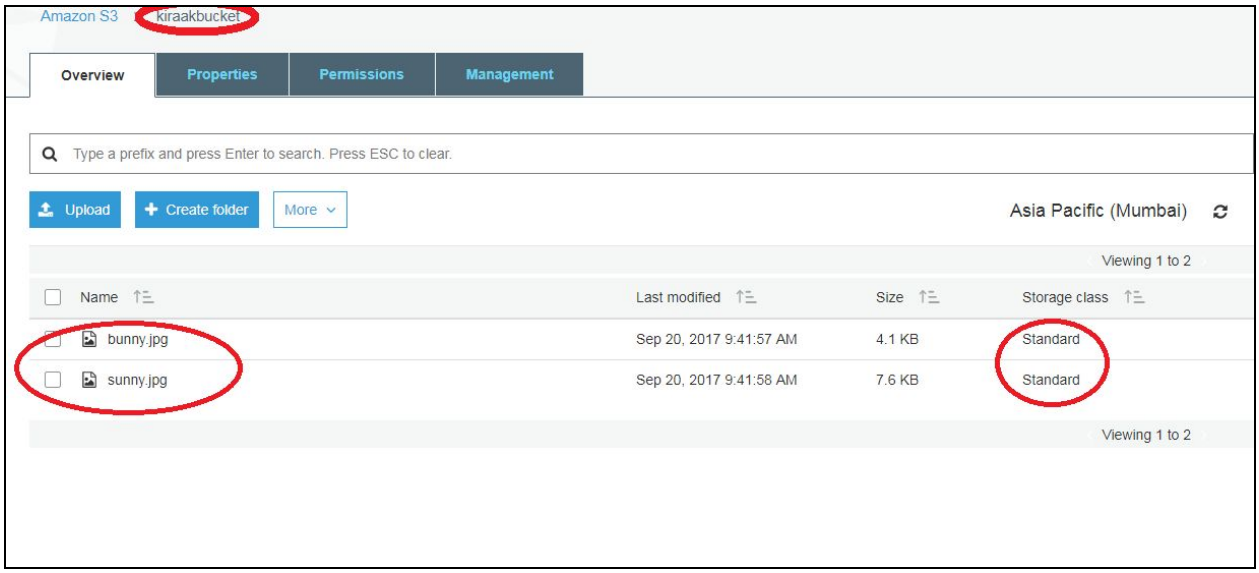

S3 Folders

S3 Folders

- For simplicity purpose, S3 supports the concept of “folders”
- This is done only as a means of grouping objects
- Amazon S3 does this by using key-name prefixes for objects

NOTE : Amazon S3 has a flat structure, there is no hierarchy like you would see in typical File System

Lab:

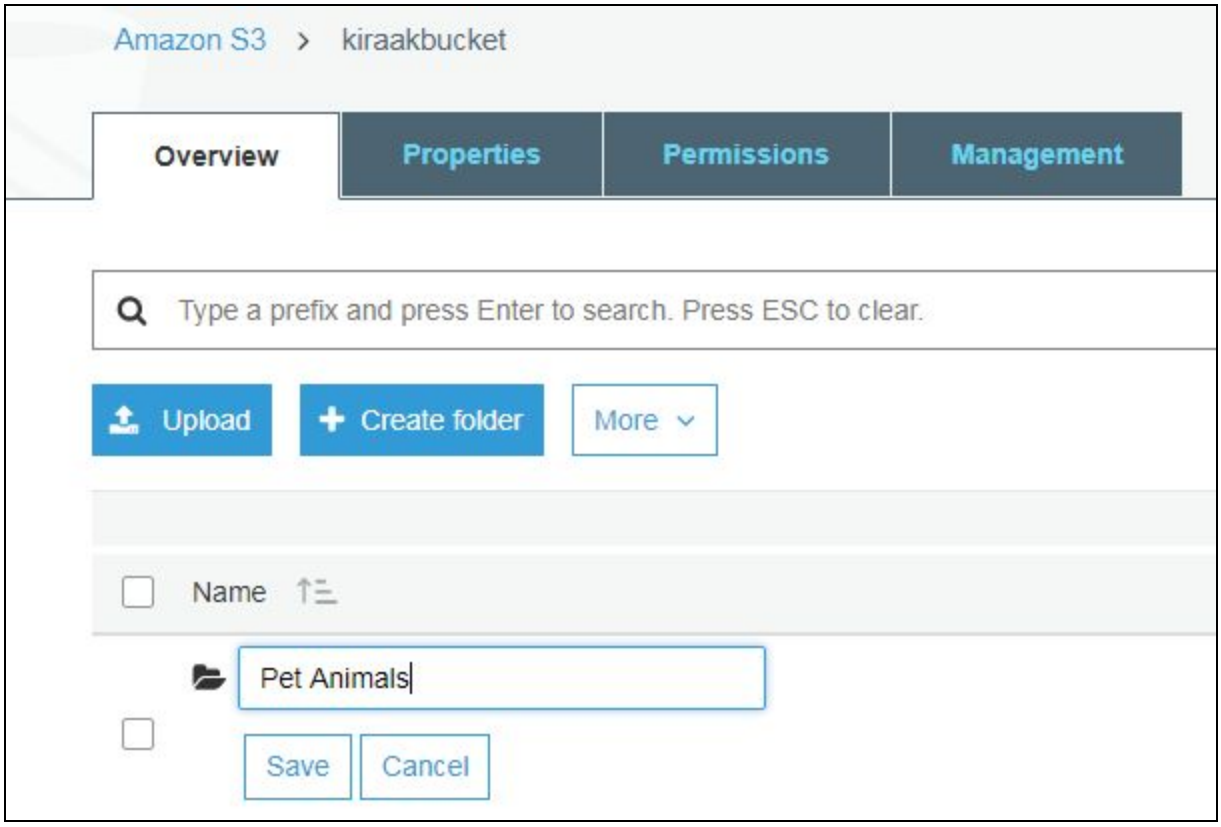

But this is just done visually. At the backend for developers purpose, it is a flat file architecture. If i upload any file in this folder then it is technically it is exactly at the same level as the other files in bucket. It is basically attaching key-name prefixes

S3 Security

Till now we have seen Bucket is created, and upload an object into this bucket. Now time to see the permissions

S3 Permissions

- All buckets and objects are private (not public) by default - only the resource owner has access.
- The resource owner can grant access to the resource (bucket/objects) through “S3 resource based policies” OR access can be granted via traditional IAM user policy

S3 Resource based policies are:

- **Bucket Policies:**

- It is the policies that are attached only to the S3 bucket (not an IAM user)
- The permissions are applied to all objects in the bucket
- The policy specifies what actions are allowed or denied for a particular user at that bucket

- **S3 ACL based policies**

- Grant access to the users in other AWS accounts or to the public
- Both buckets and objects has ACLs
- Object ACL allow us to share an S3 object with the public via a URL link

LAB 3: Try to Access that objects from browser(Permission)

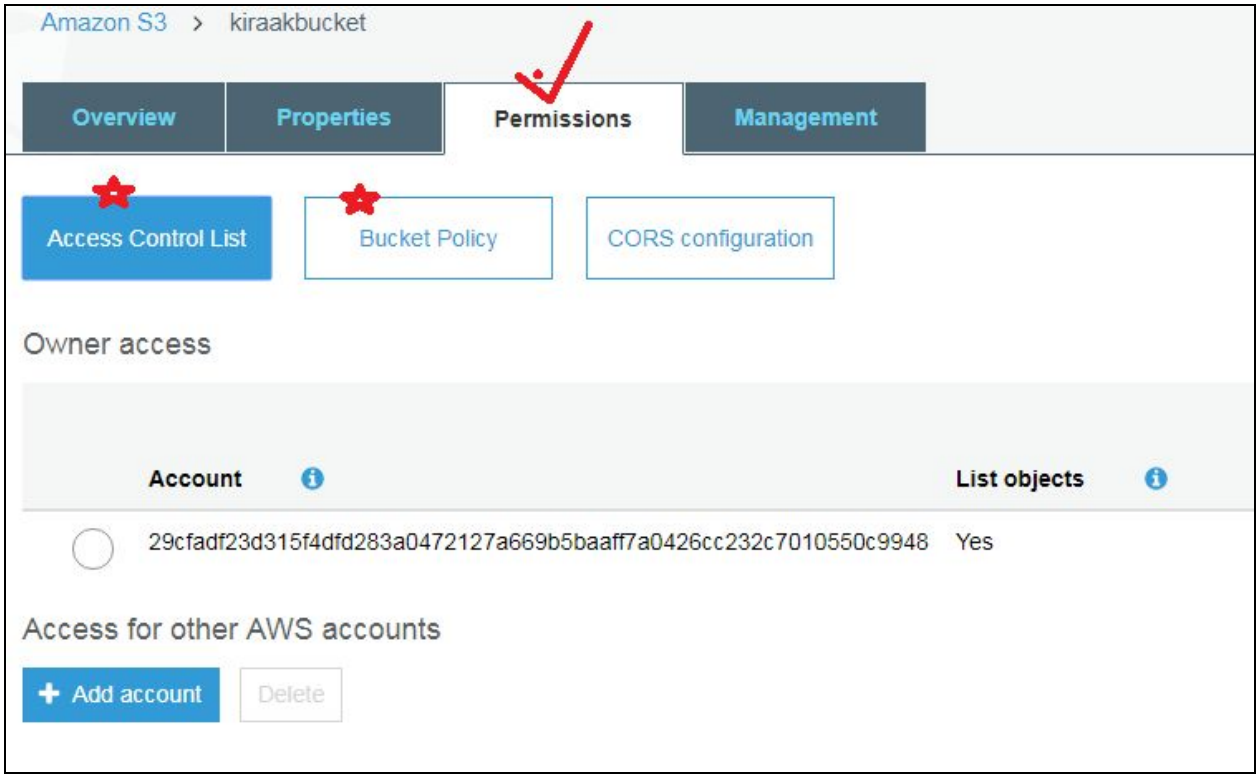

Now go the objects and click on link
By default u can't open in public so give the permission as "make public" then u can access from anywhere.

Also take the access back.

S3 Storage Class

S3 Storage Class:

- A storage class represents the “classification” assigned to each Object in S3. There are 4 types :

1. Standard
2. Reduced Redundancy Storage (RRS)
3. Infrequently Access (S3-IA)
4. Glacier

- Each storage class has varying attributes that dictate things like
- Cost
- Object Availability
- Object Durability
- Frequency of access (to the object)

STANDARD

→ Designed for general, all-purpose storage
→ It is the default storage option
→ 99.999999999% object durability(eleven 9’s)
→ 99.99% of object availability
→ It is the most expensive storage class

Reduced Redundancy (RRS)

→ Designed for non-critical, reproducible objects
→ 99.99% object durability
→ 99.99% object availability
→ it is less expensive than the standard

Infrequent Access (S3-IA)

→ Designed for objects that you do not frequently access, but must be immediately available when accessed
→ 99.999999999% object durability
→ 99.90% object availability
→ It is less expensive than the Standard/RRS

Glacier

- Designed for long term archival storage (not to be used for backup)
- May take several hours for objects stored in glacier to be retrieved
- 99.999999999% of object durability
- It is cheapest S3 storage class among all (very low cost)

Storage classes features in table

	Standard	Standard - Infrequent Access	Reduced Redundancy Storage
Durability	99.999999999%	99.999999999%	99.99%
Availability	99.99%	99.9%	99.99%
Concurrent facility fault tolerance	2	2	1
SSL support	Yes	Yes	Yes
First byte latency	Milliseconds	Milliseconds	Milliseconds
Lifecycle Management Policies	Yes	Yes	Yes

LAB 4: Selecting Storage Type

Go to the object under bucket and then select properties tab
Select storage class

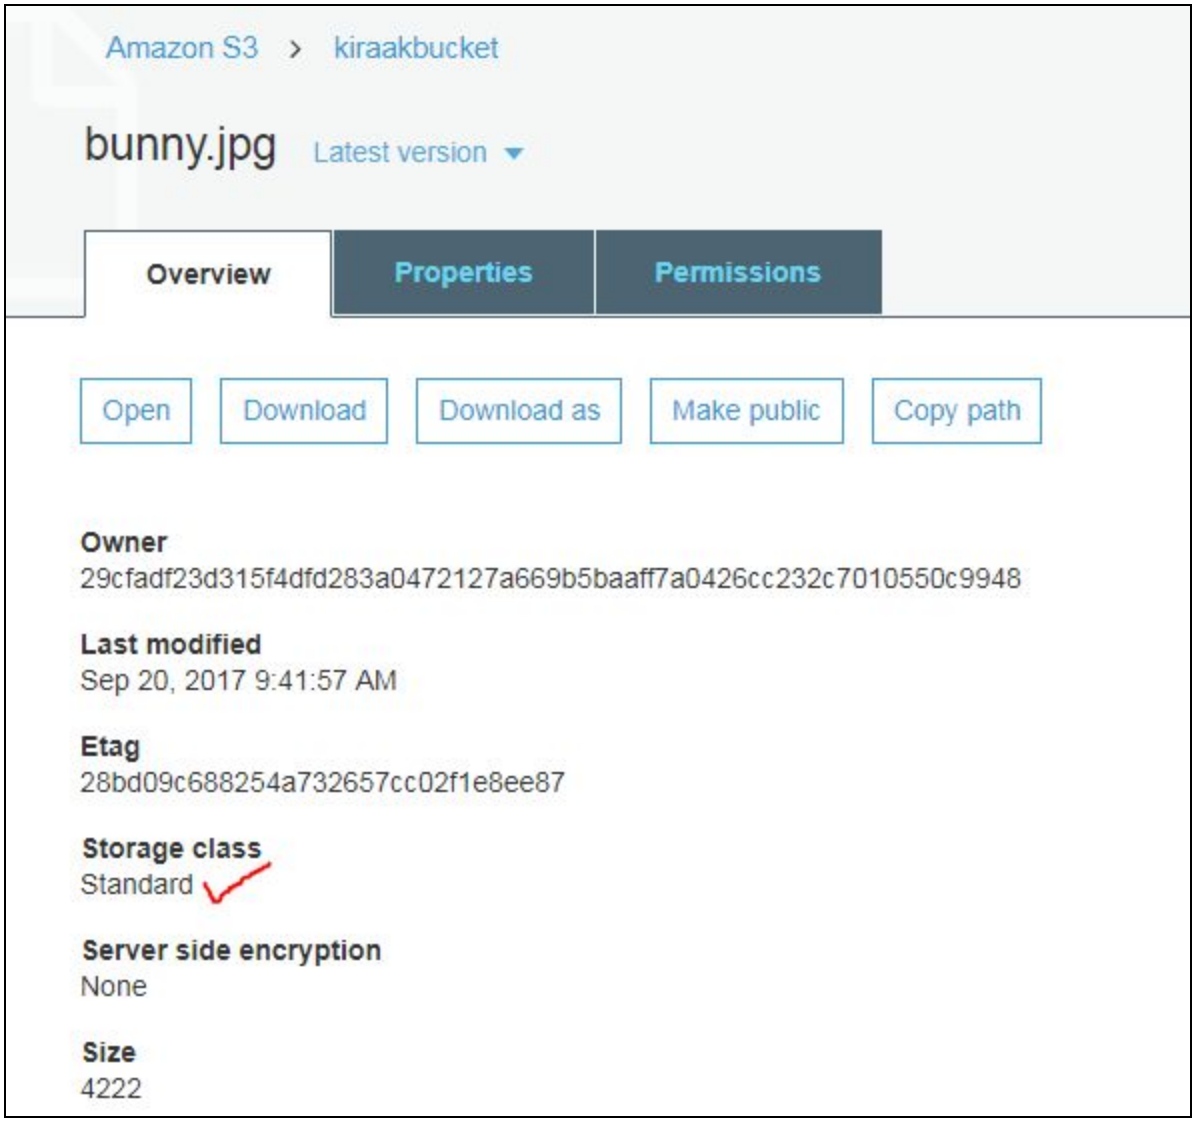

Click on Properties

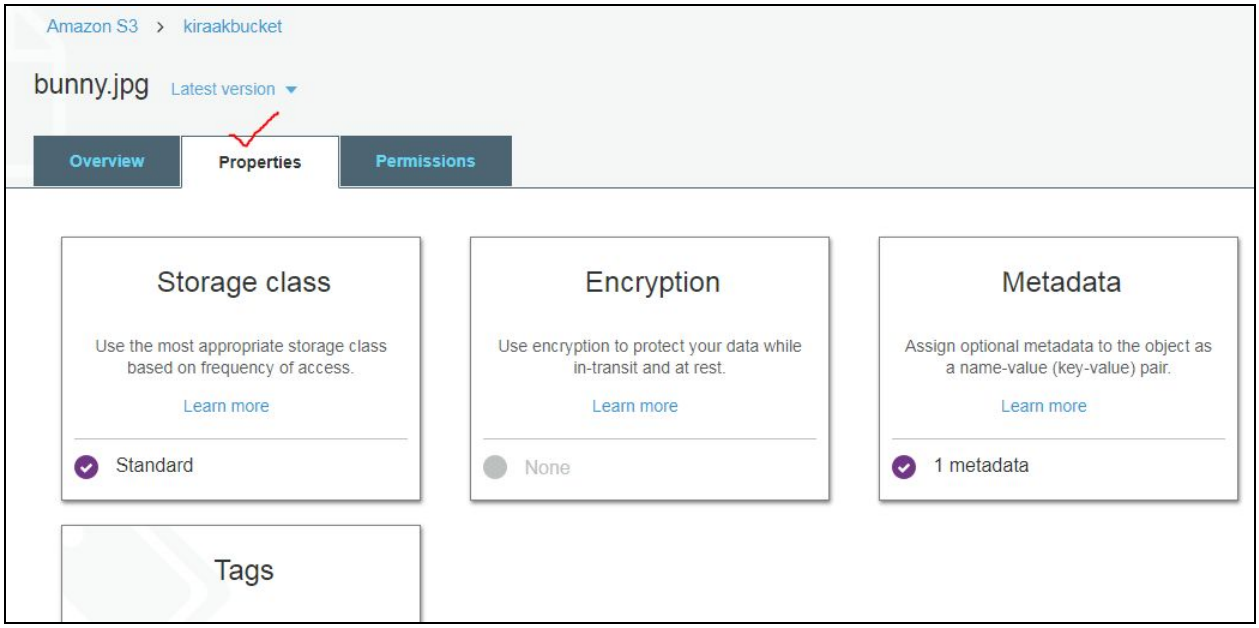

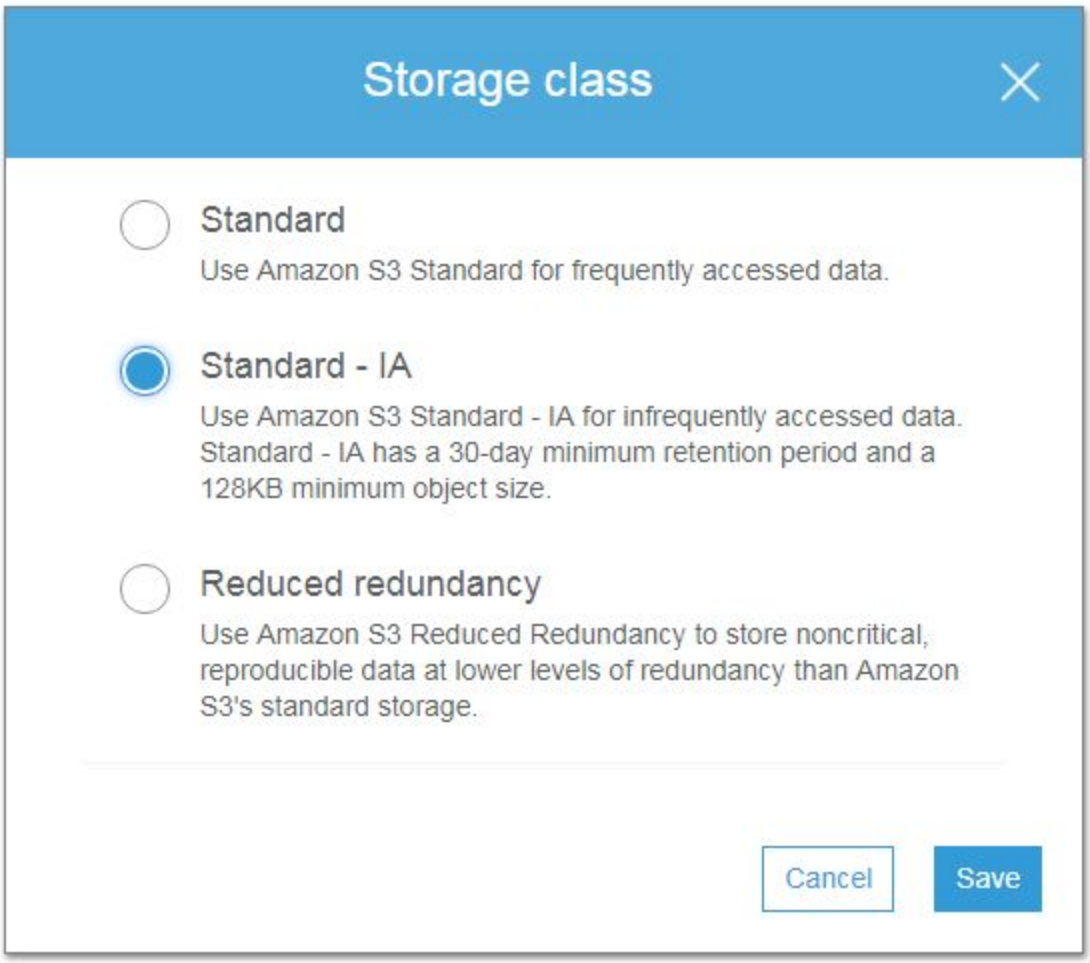

Click on save

You can verify

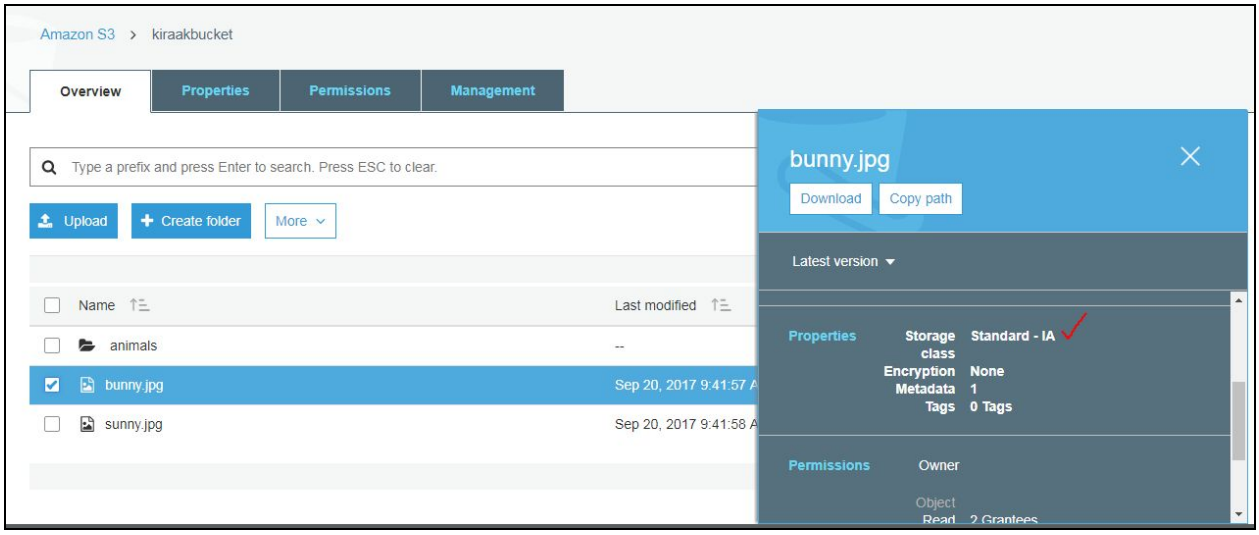

Object Versioning

- S3 versioning is a feature to manage and store all old/new/deleted versions of an object.
- By default, versioning is disabled on all buckets/objects
- Once versioning is enabled, you can only suspend versioning. It cannot be fully disabled
- Suspending versioning only prevents new versions from being created. All objects with existing versions will maintain their older versions
- Versioning can only be set on the bucket level and applies to ALL objects in the bucket
- Lifecycle policies can be applied to specific versions of an object
- Versioning and lifecycle policies can both be enabled on a bucket at the same time
- Versioning can be used with lifecycle policies to create a great archiving and backup solutions in s3

Lab 5: Versioning Exercise

Go to the bucket and upload one file

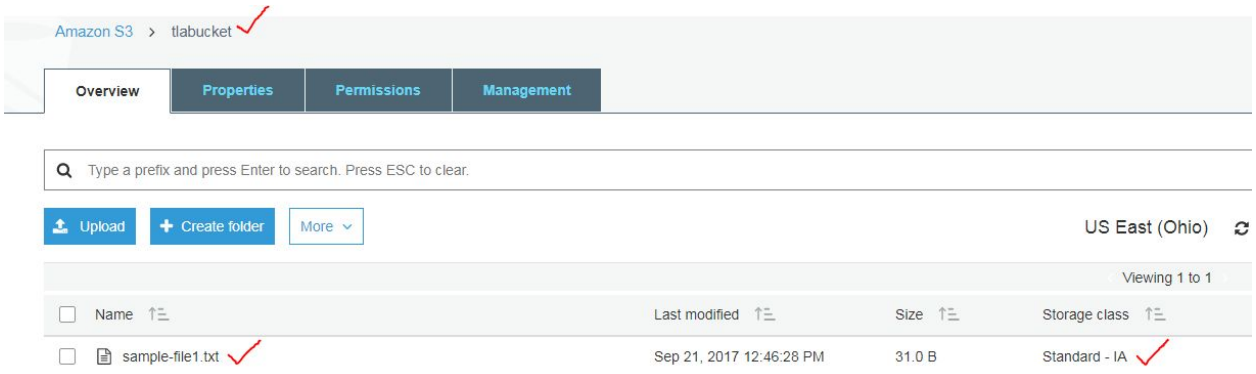

Now enable the versioning in the bucket and then change the file

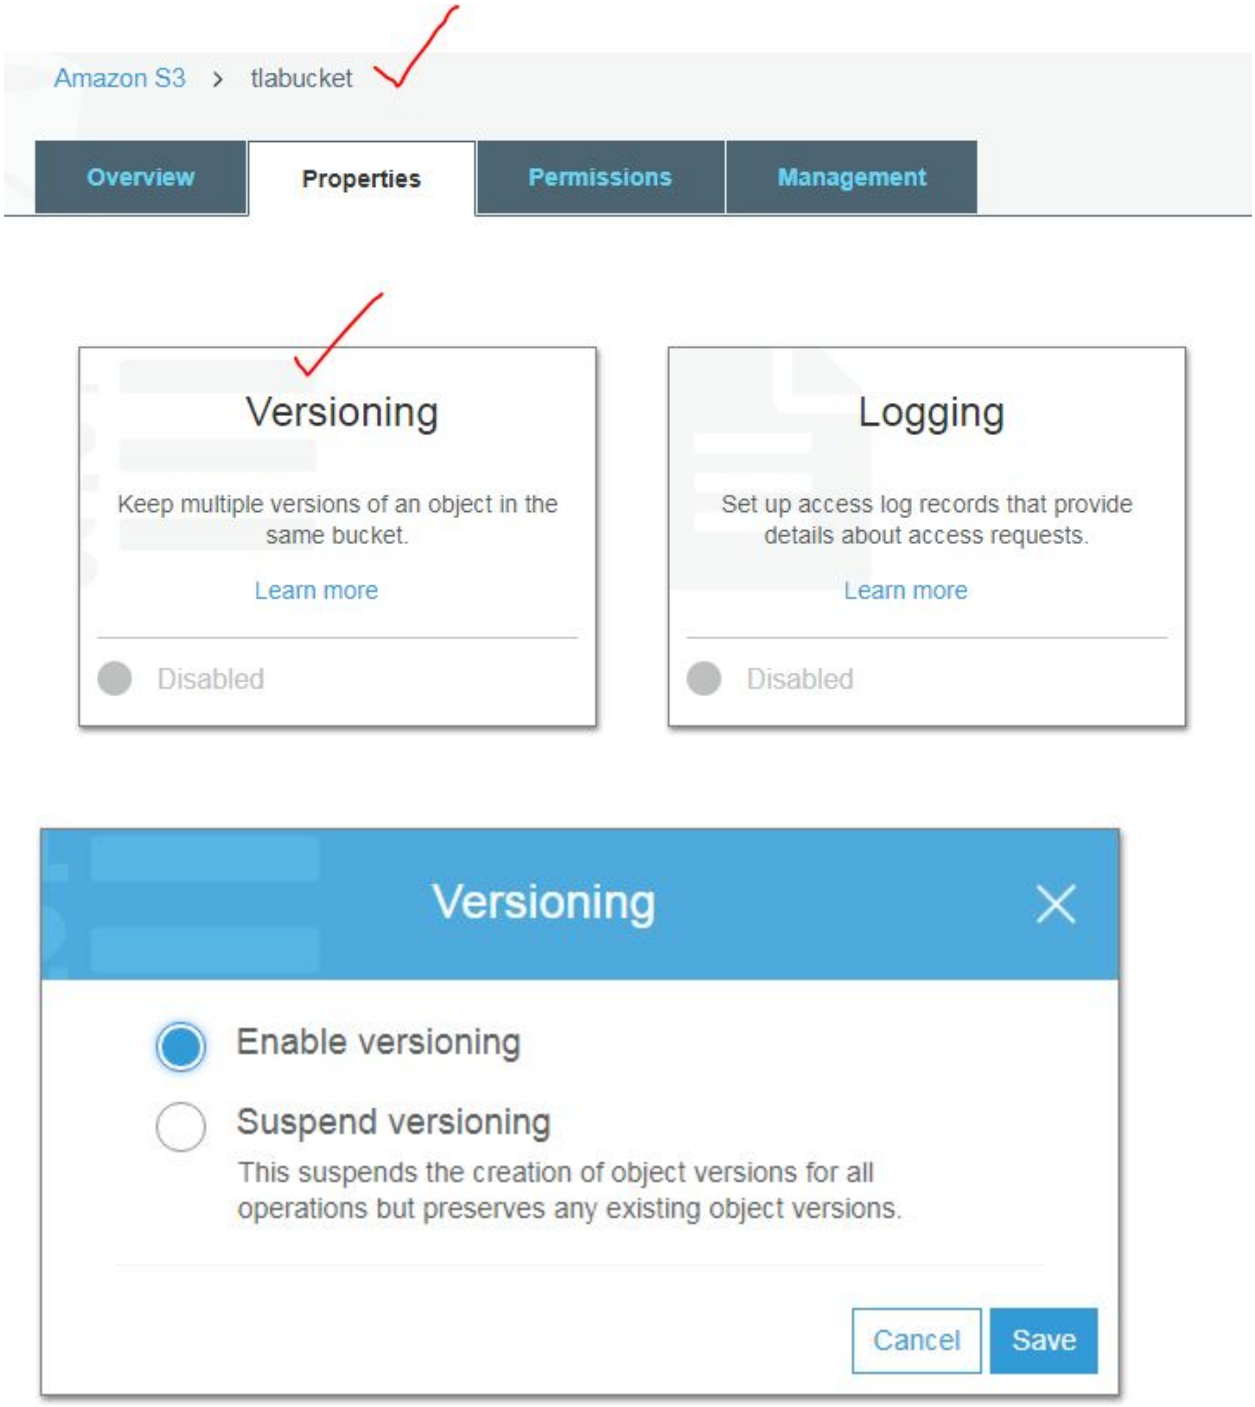
Amazon S3 > tlabucket
Overview
Properties
Permissions
Management
Versioning
Keep multiple versions of an object in the same bucket.
Learn more
Disabled
Logging
Set up access log records that provide details about access requests.
Learn more
Disabled
Versioning
Enable versioning
Suspend versioning
This suspends the creation of object versions for all operations but preserves any existing object versions.
Cancel
Save

Now change the file content and again upload same file in bucket

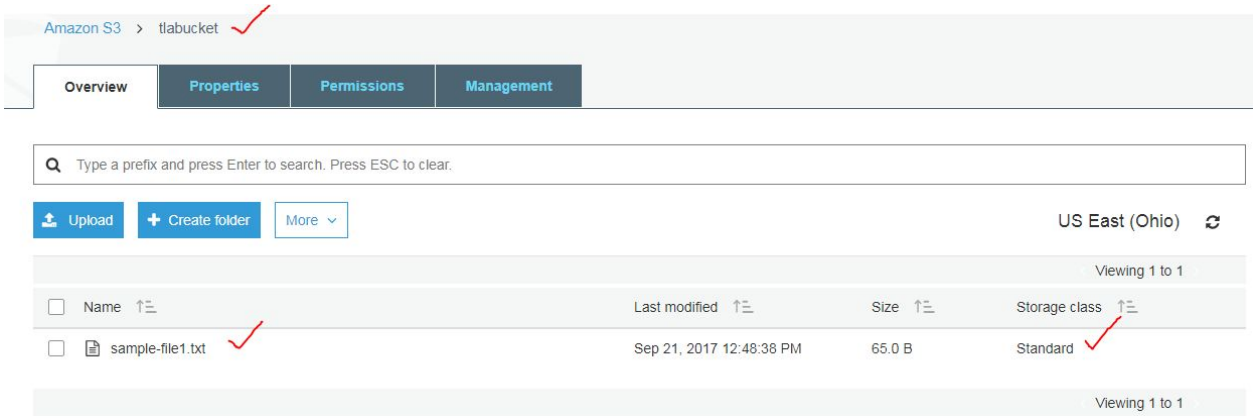

This is the new file

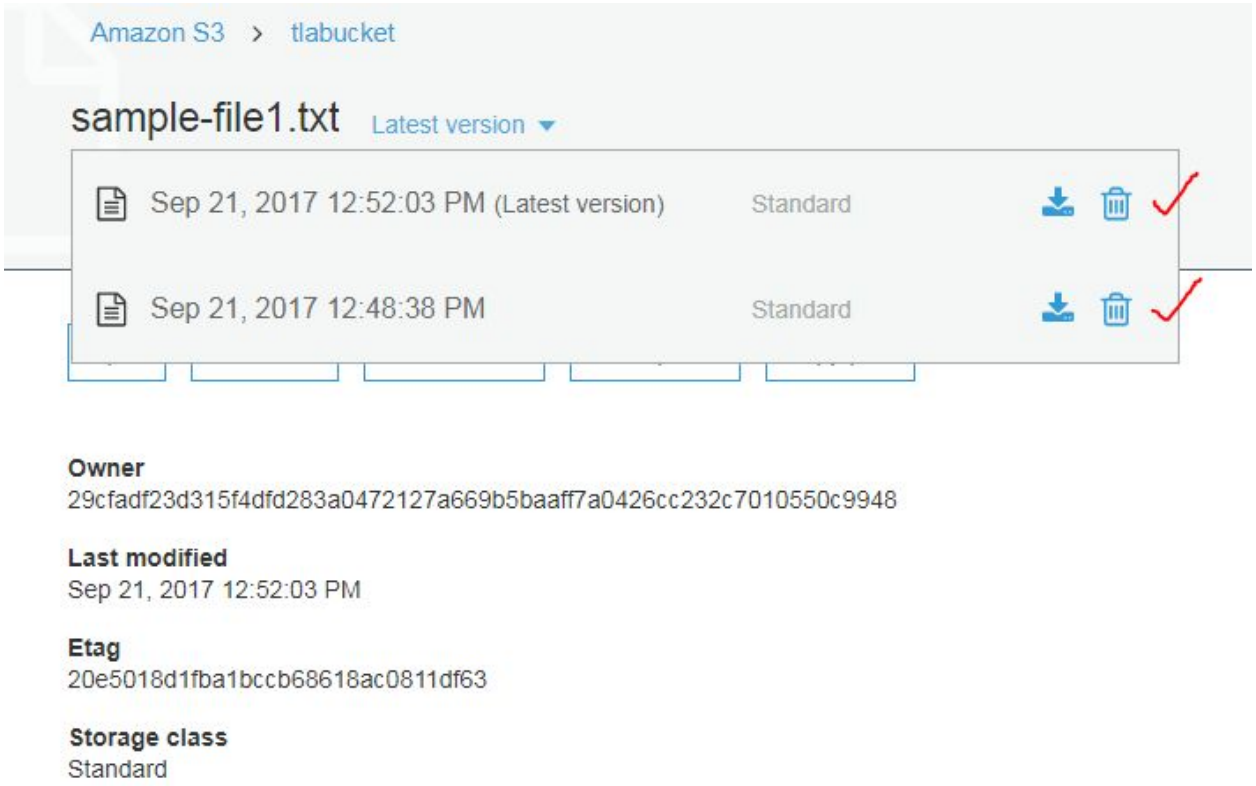

You can also see the various versions of the existing files.

Lifecycle Policies

- An object lifecycle policy is a **set of rules that automate** the migration of an object's storage class to a different storage class (or deletion), based on specified intervals
- By default, lifecycle policies are disabled on a bucket/objects
- These policies are customizable to meet your company's data retention policies
- Great for automating the management of object storage and to be more cost efficient
- Can be used with versioning to create a great archiving and backup solutions in s3

Scenario:

1) I have a work file that I am going to access every day for the next 30 days
2) After 30 days, I may only need to access that file once a week for the next 60 days
3) After which (90 days total) I will probably never access the file again but want to keep it just in case.

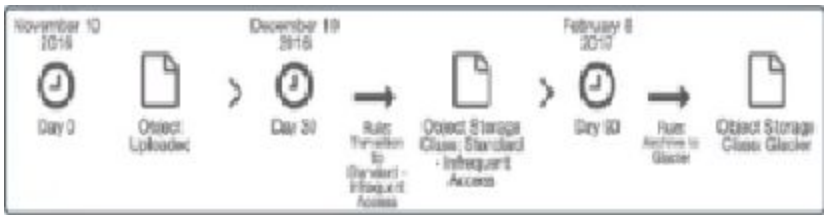

LAB 6: Setup the Lifecycle policy to the bucket as per req

Go to bucket and then click on management tab, then click on lifecycle

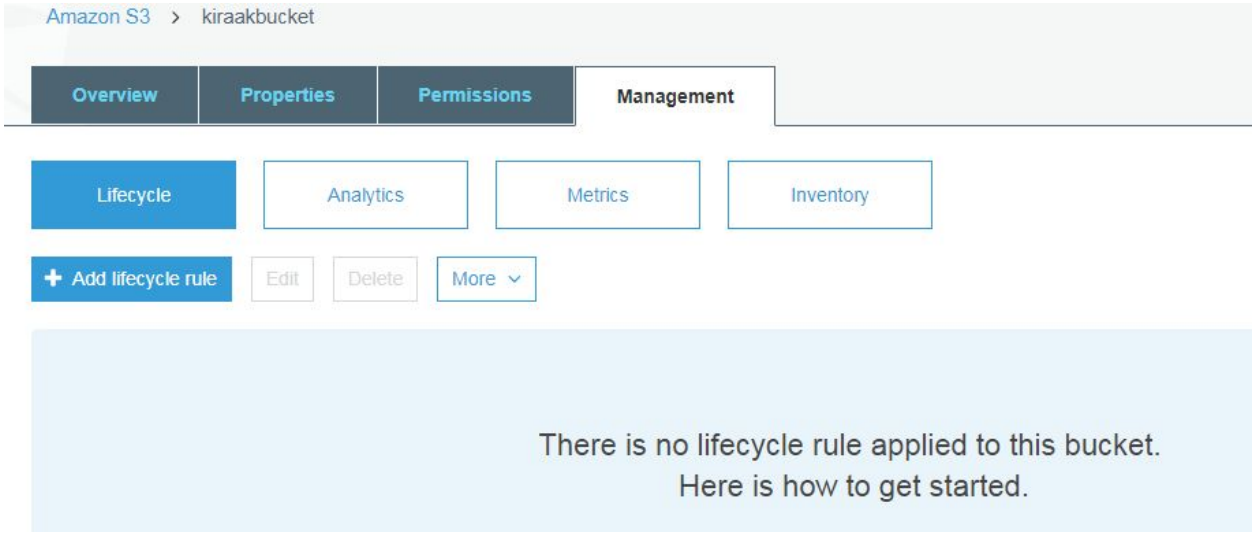

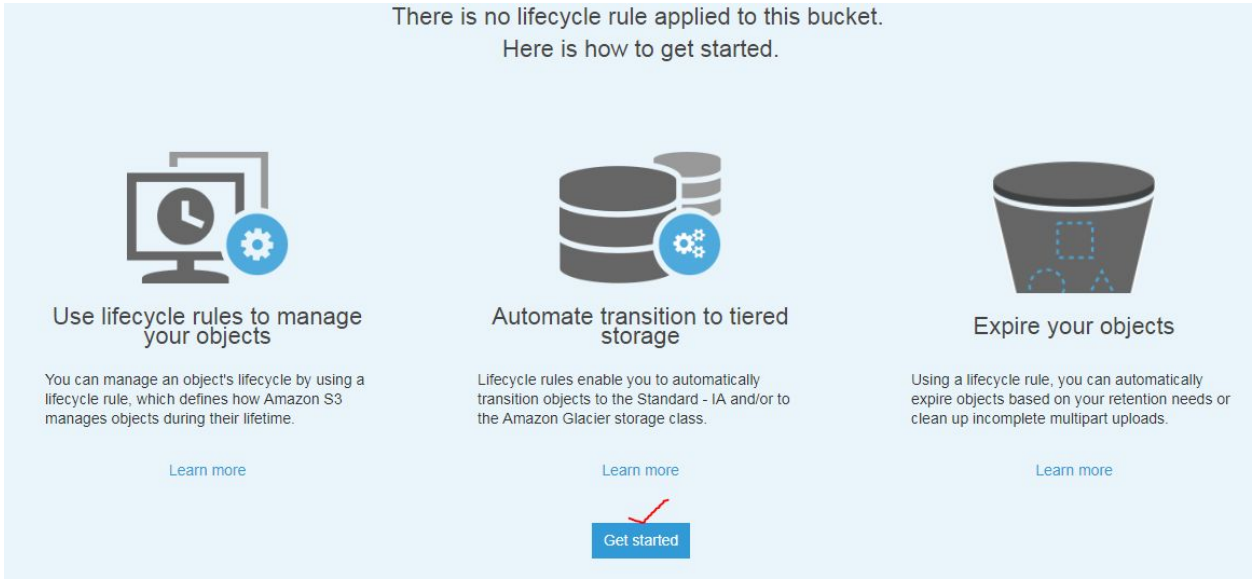

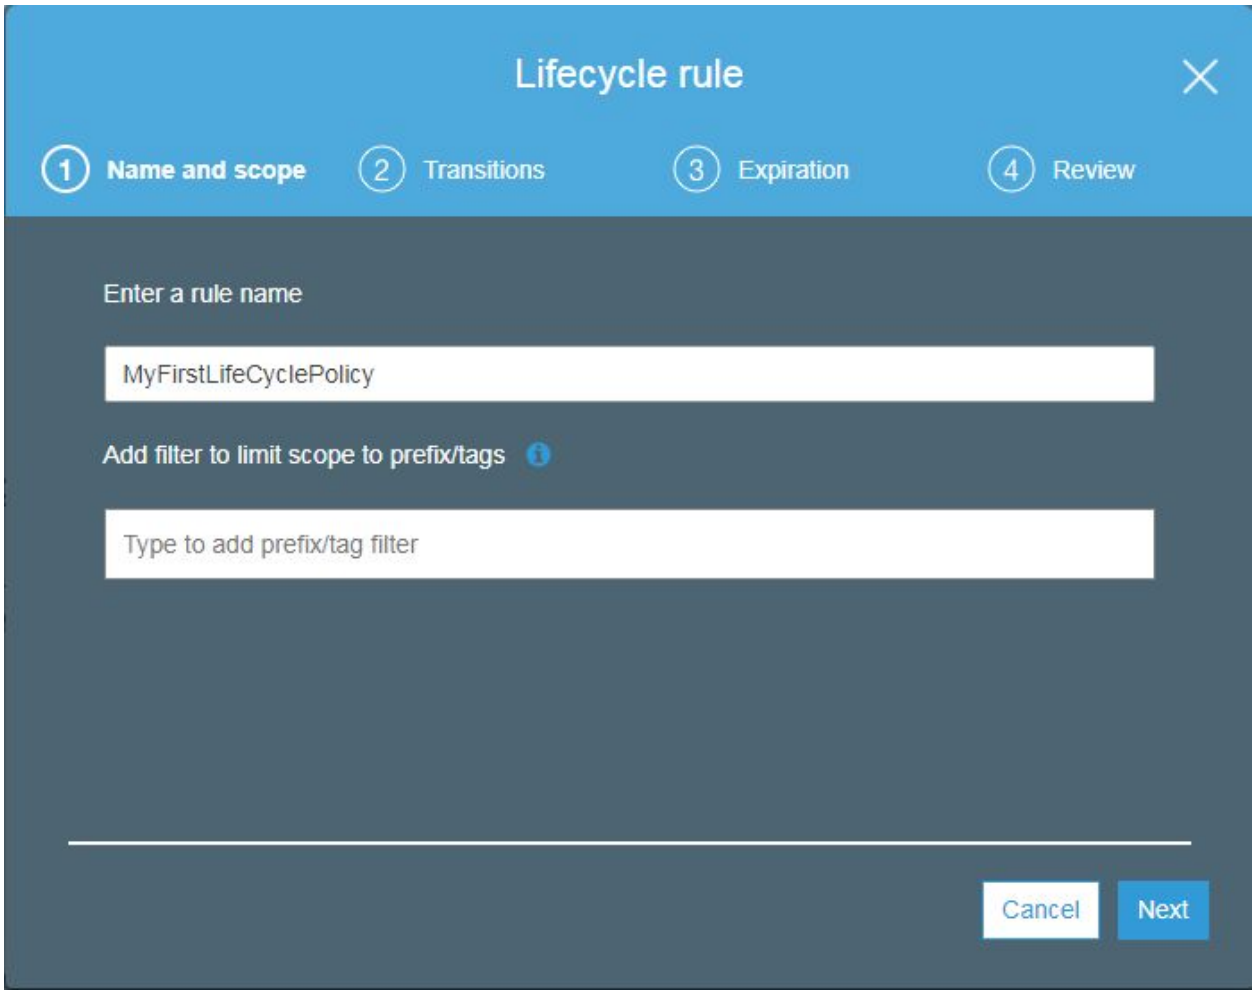
Lifecycle rule
1 Name and scope
2 Transitions
3 Expiration
4 Review
Enter a rule name
MyFirstLifeCyclePolicy
Add filter to limit scope to prefix/tags
Type to add prefix/tag filter
Cancel
Next

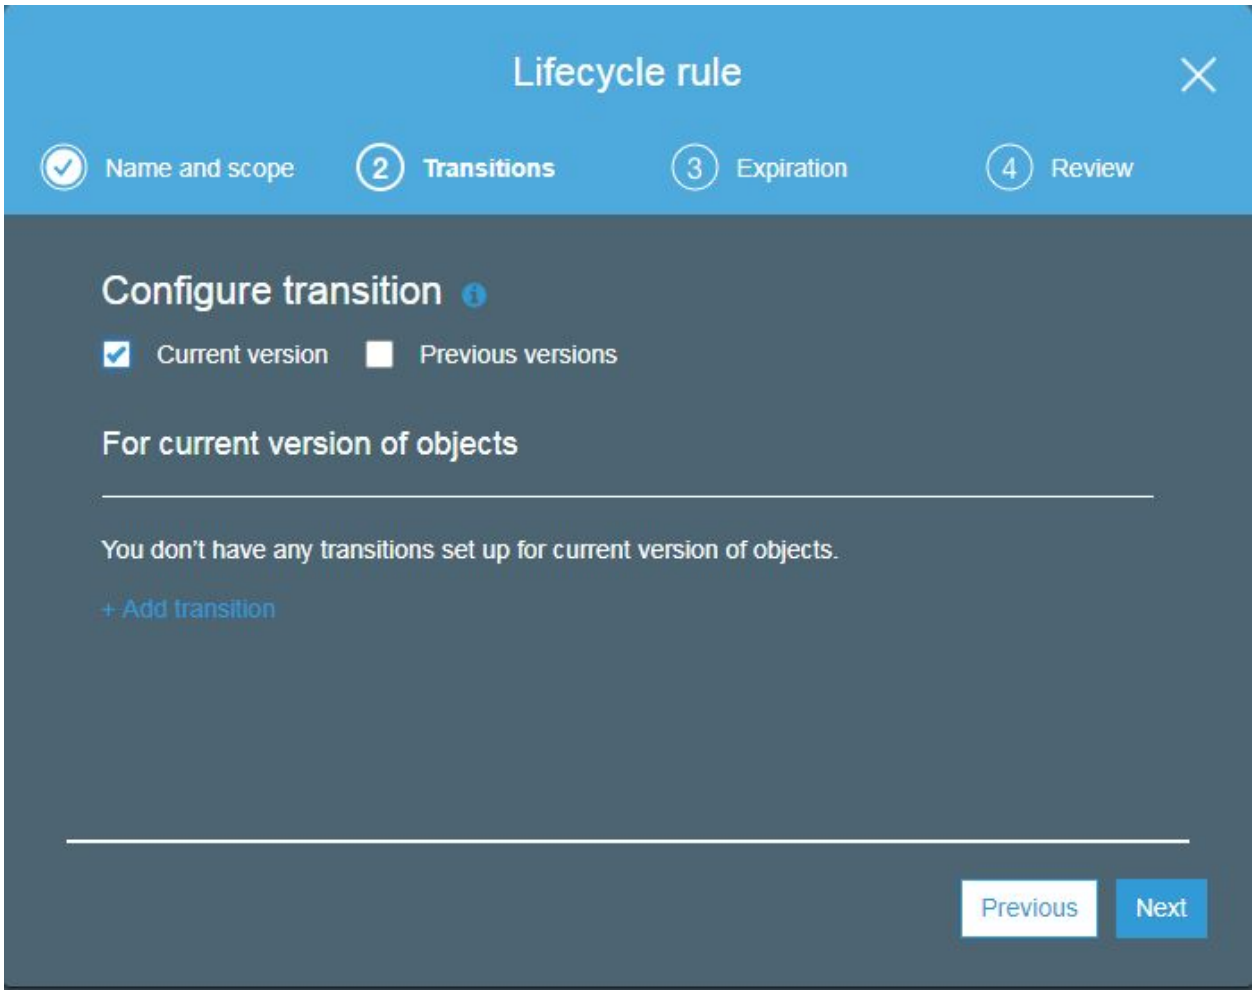
Lifecycle rule
Name and scope
2
Transitions
3
Expiration
4
Review
Configure transition
Current version
Previous versions
For current version of objects
You don't have any transitions set up for current version of objects.
+ Add transition
Previous
Next

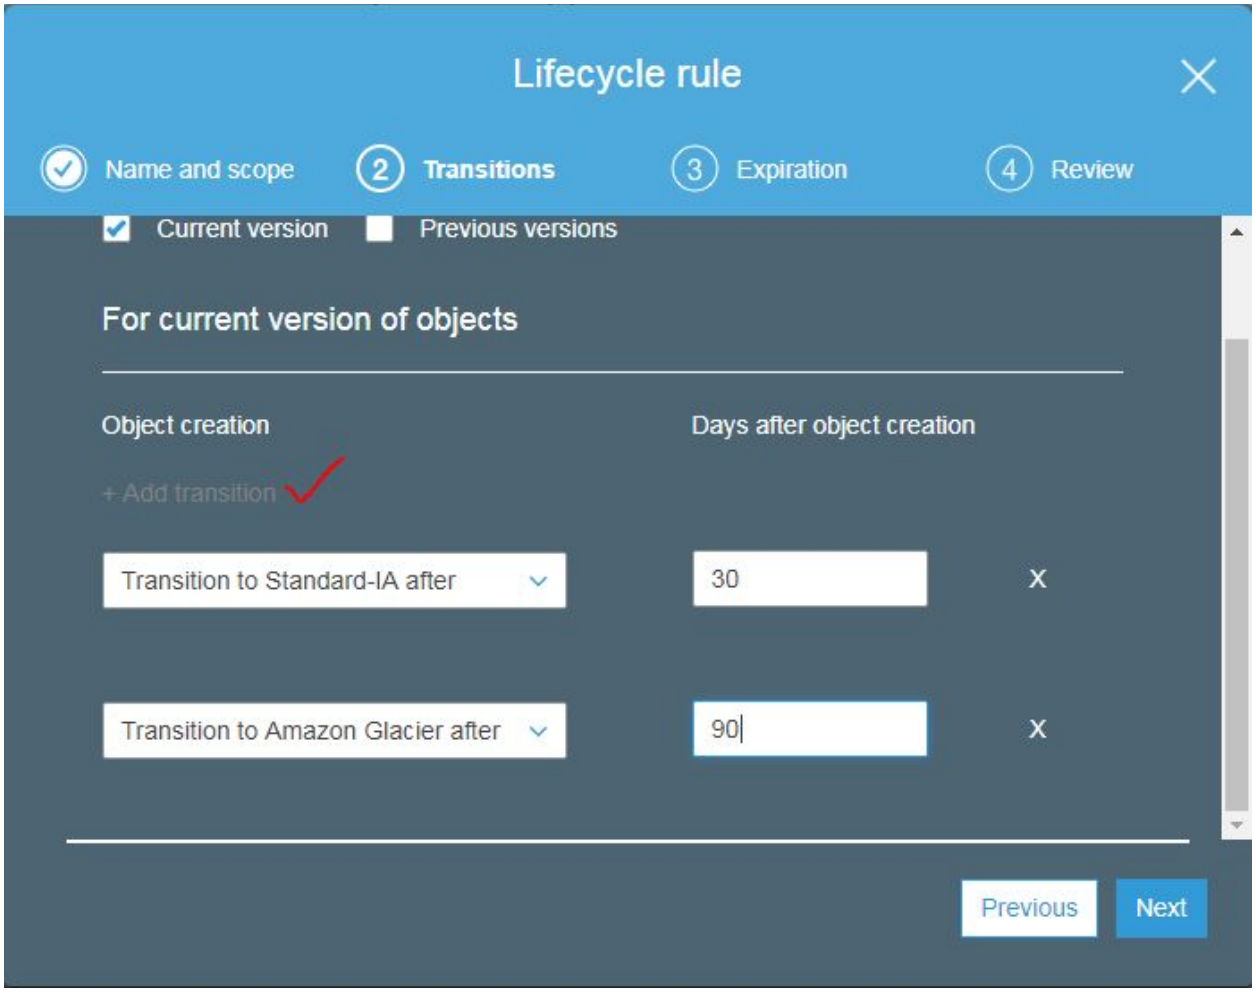
Lifecycle rule
Name and scope
2 Transitions
3 Expiration
4 Review
Current version
Previous versions
For current version of objects
Object creation
Days after object creation
+ Add transition
Transition to Standard-IA after
30
X
Transition to Amazon Glacier after
90
X
Previous
Next

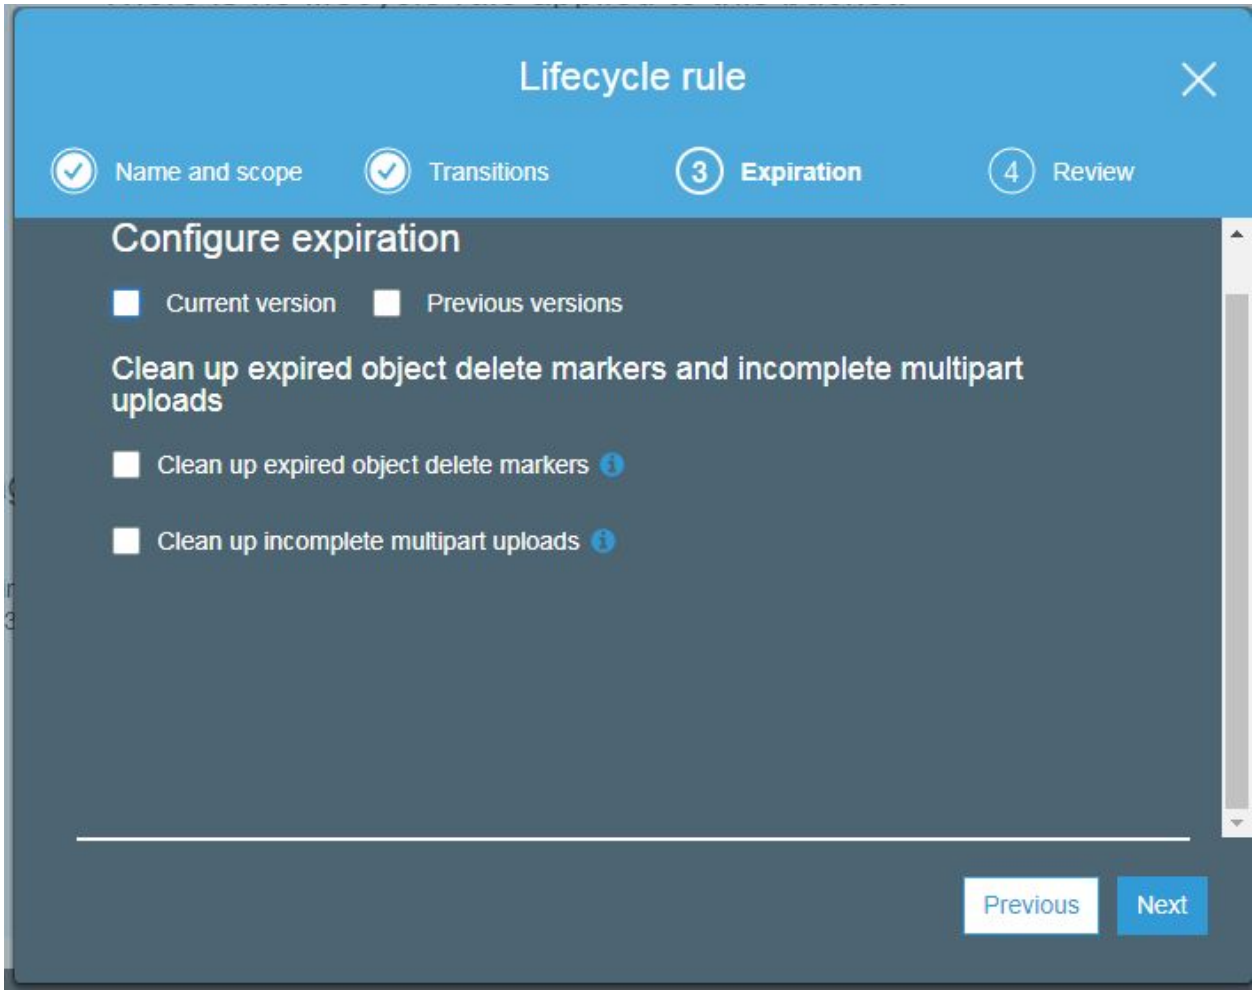
Lifecycle rule
Name and scope
Transitions
3
Expiration
4
Review
Configure expiration
Current version
Previous versions
Clean up expired object delete markers and incomplete multipart uploads
Clean up expired object delete markers
Clean up incomplete multipart uploads
Previous
Next

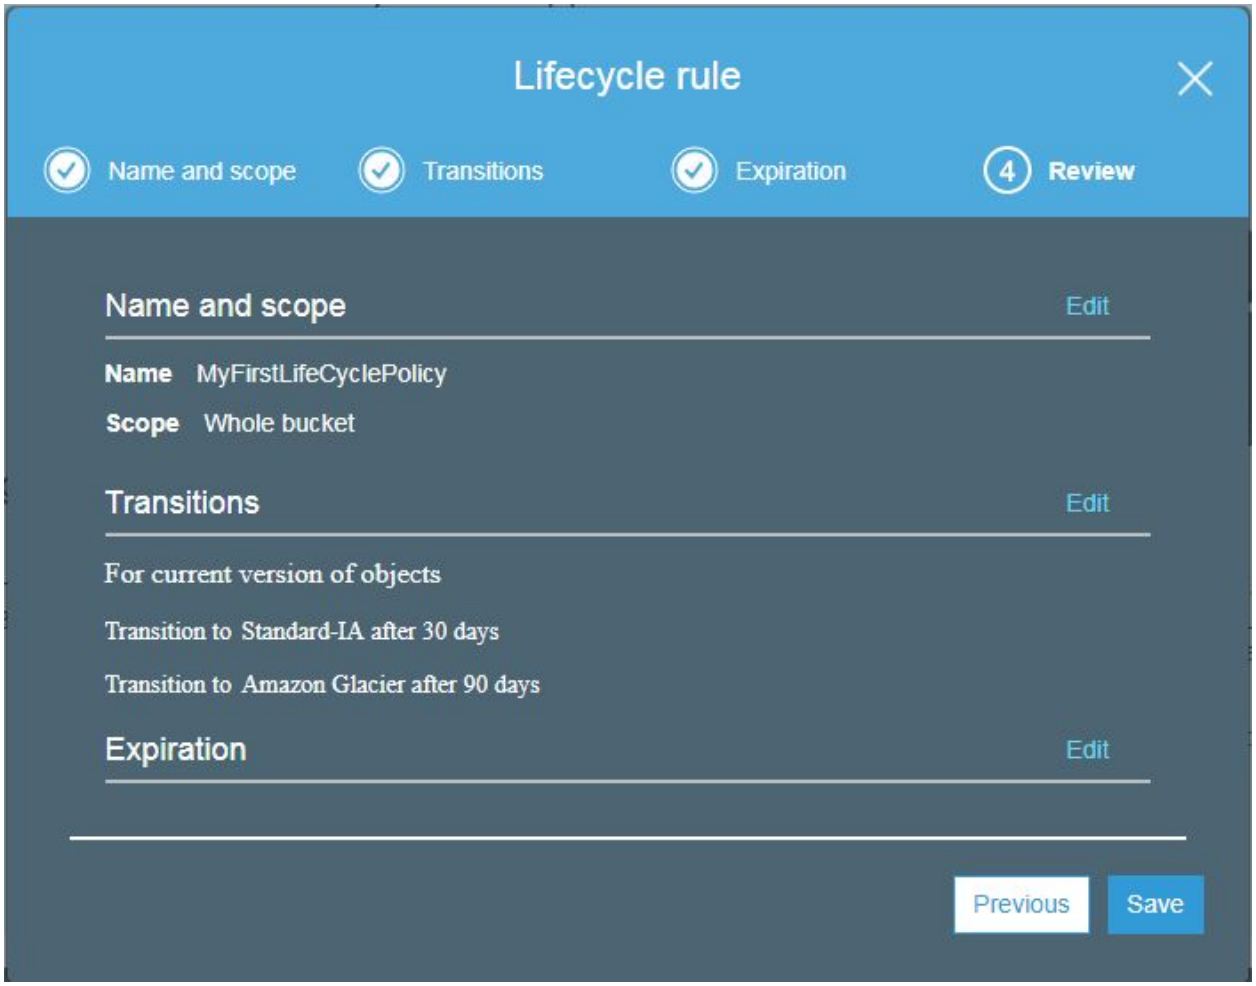
Lifecycle rule
Name and scope
Transitions
Expiration
4 Review
Name and scope
Edit
Name MyFirstLifeCyclePolicy
Scope Whole bucket
Transitions
Edit
For current version of objects
Transition to Standard-IA after 30 days
Transition to Amazon Glacier after 90 days
Expiration
Edit
Previous
Save

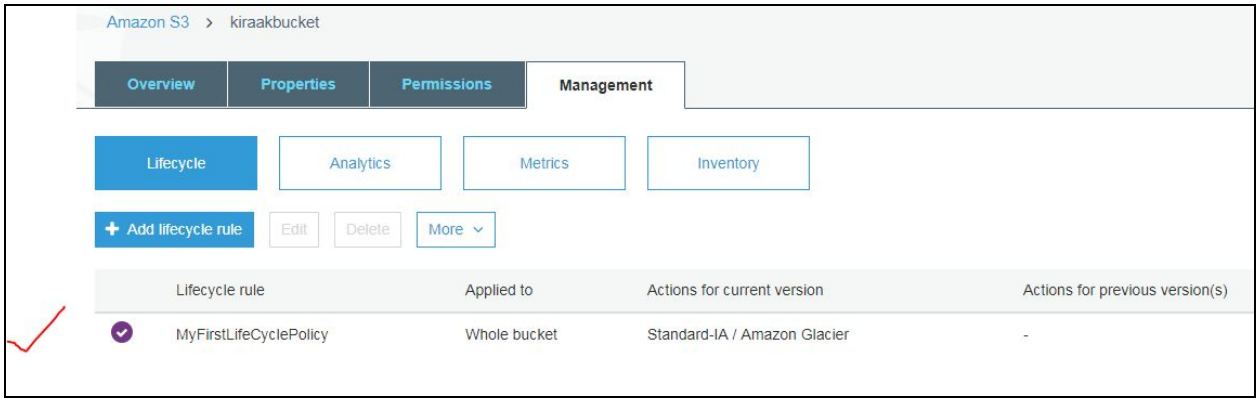
Amazon S3 › kiraakbucket
Overview
Properties
Permissions
Management
Lifecycle
Analytics
Metrics
Inventory
+ Add lifecycle rule
Edit
Delete
More
Lifecycle rule
Applied to
Actions for current version
Actions for previous version(s)
MyFirstLifeCyclePolicy
Whole bucket
Standard-IA / Amazon Glacier
-

S3 Static Web Hosting

- AWS S3 provides an option for a low-cost, highly reliable web hosting service for a static websites (content that does not change frequently)
- When enabled, static web hosting will provide you with a unique endpoint (url) that you can point to any properly formatted file stored in an s3 bucket. Supported formats include:

- HTML
- CSS
- JavaScript

- AWS Route53 also helps to map human readable domain names to static web hosting buckets, which are ideal for DNS failover solutions

LAB 7: Web Hosting the static site using S3

Create 2 file by name index.html and error.html and upload into bucket

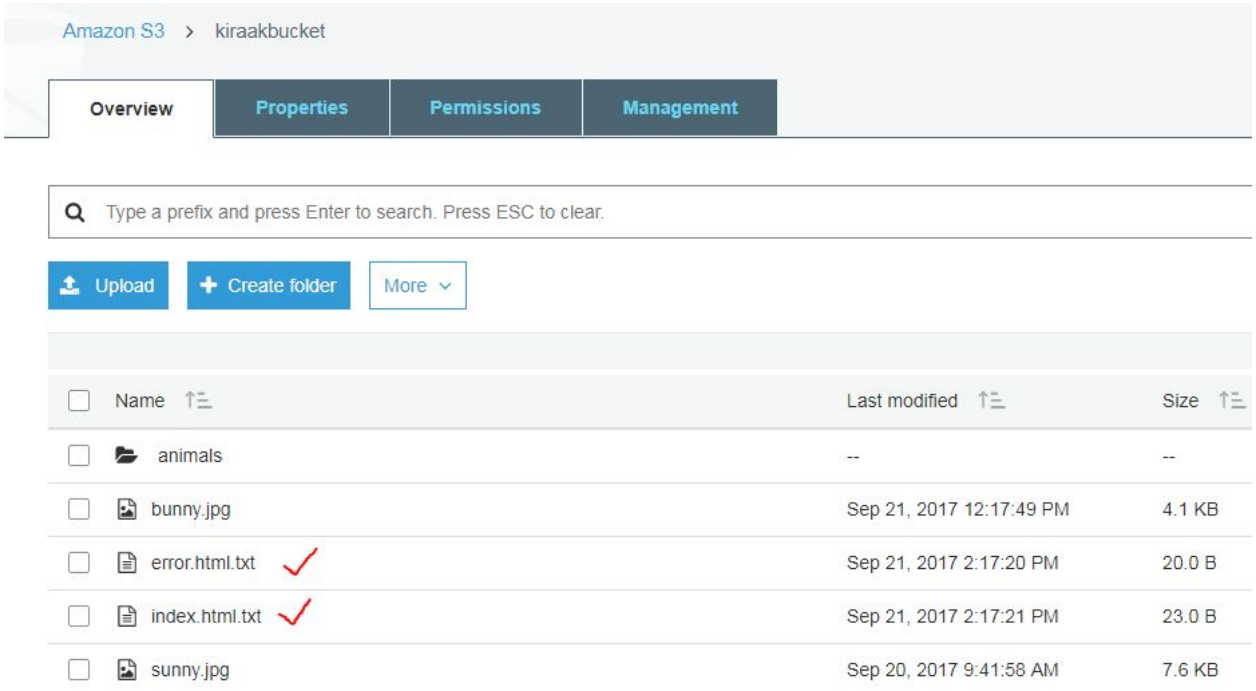

Now after that goto properties and click on static web hosting

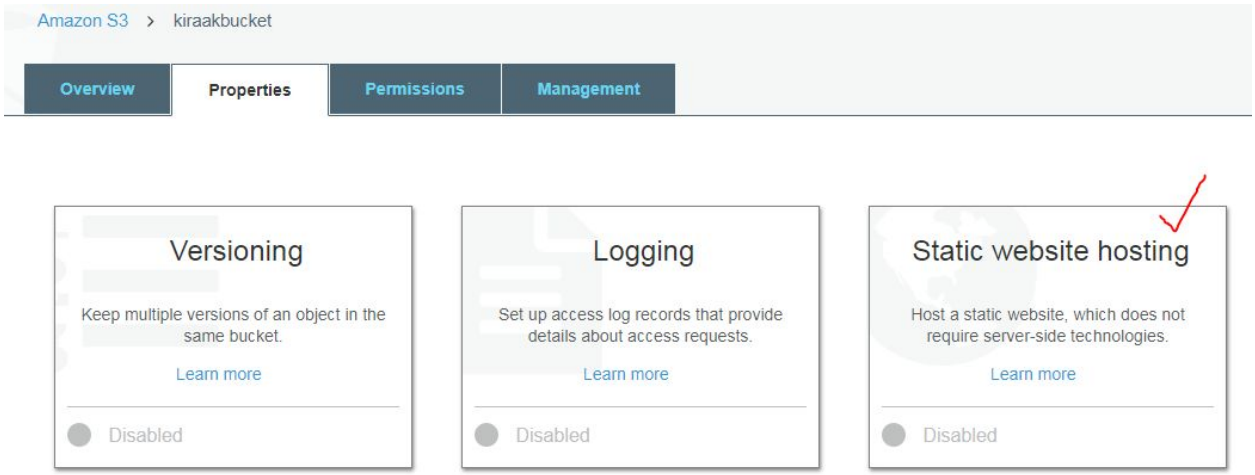
Amazon S3 › kiraakbucket
Overview
Properties
Permissions
Management
Versioning
Keep multiple versions of an object in the same bucket.
Learn more
Disabled
Logging
Set up access log records that provide details about access requests.
Learn more
Disabled
Static website hosting
Host a static website, which does not require server-side technologies.
Learn more
Disabled

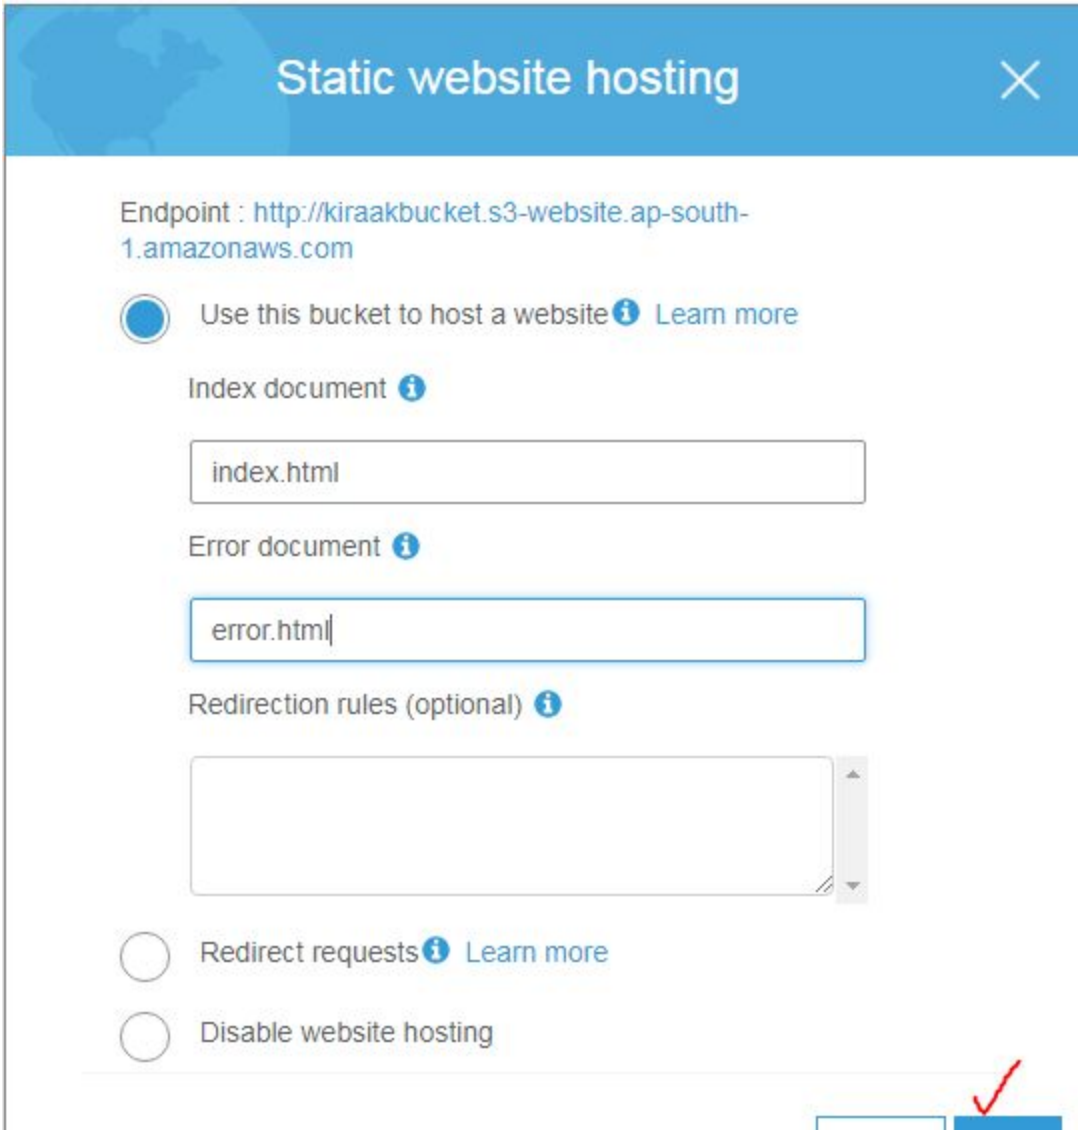
Static website hosting
Endpoint : http://kiraakbucket.s3-website.ap-south-1.amazonaws.com
Use this bucket to host a website
Learn more
Index document
index.html
Error document
error.html
Redirection rules (optional)
Redirect requests
Learn more
Disable website hosting

Now you can copy url by going inside static hosting tab

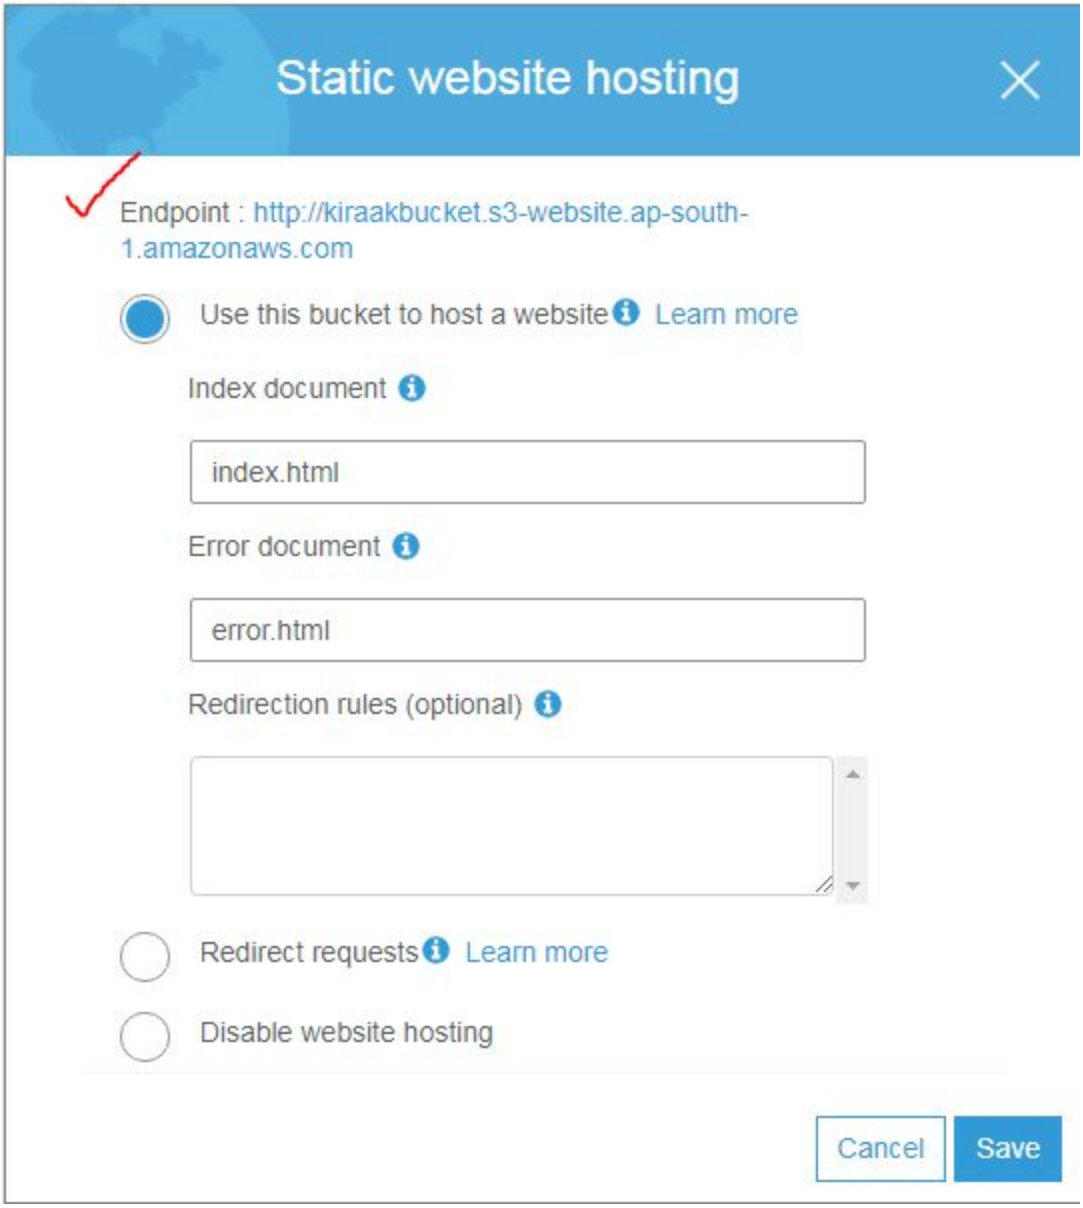

Paste that url in browser

403 Forbidden

- Code: AccessDenied
- Message: Access Denied
- RequestId: C3BB9923740704AE
- HostId: L77heHjKLCx1bwg3ydq48/+bYV4Zzv8zAYrr9pN+MWUBU5QXAE96Wfn8C/LWSYxjqD411b1D0Yc=

An Error Occurred While Attempting to Retrieve a Custom Error Document

- Code: AccessDenied
- Message: Access Denied

You get the error as there is no proper permission to public

GO to the objects and give the public permission

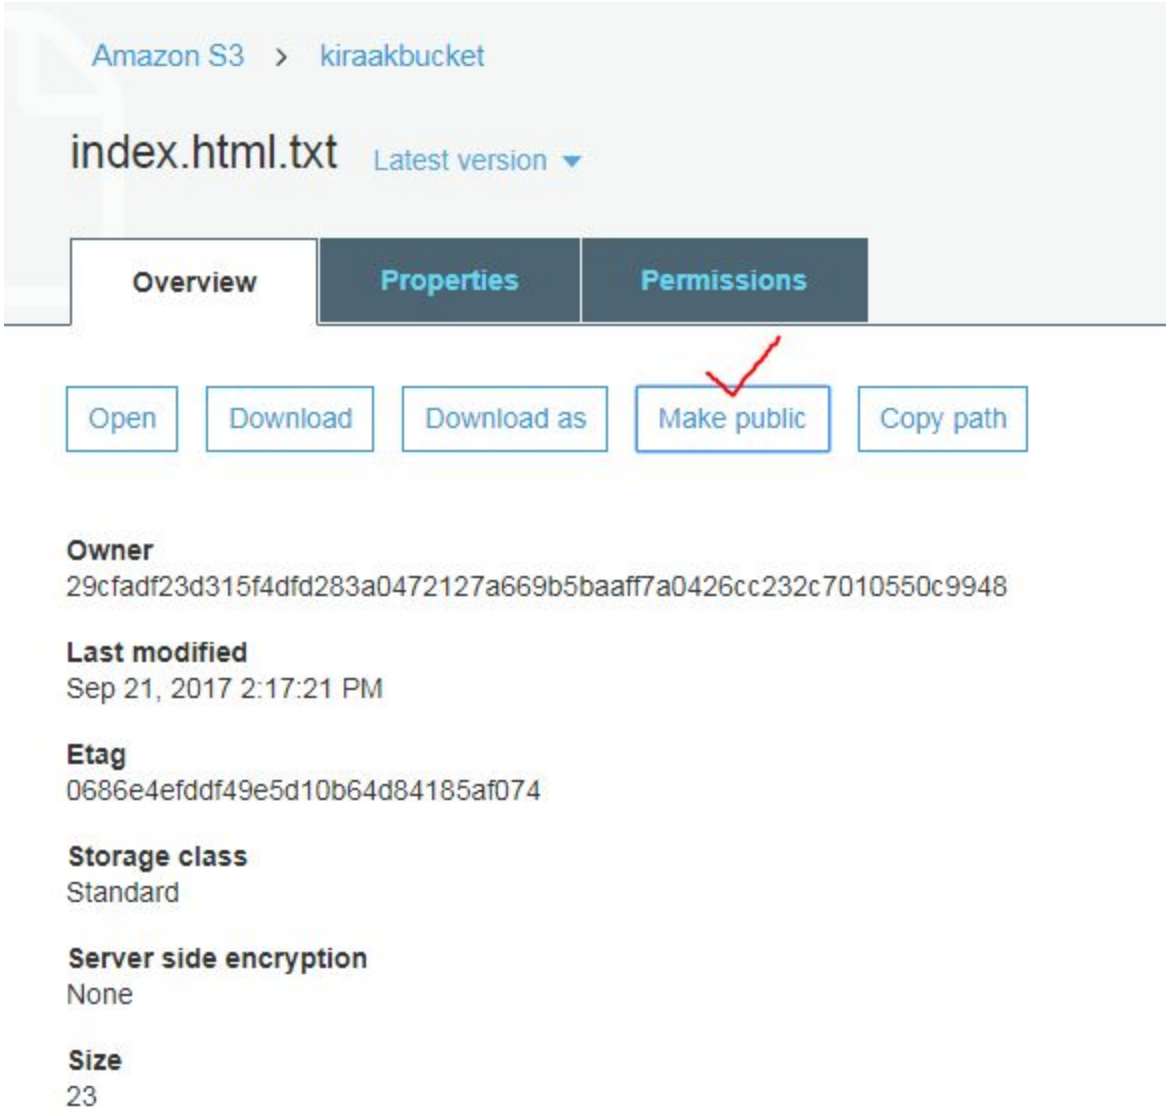

Same for error.html file too

CORS

- CORS is a method of allowing a web application located in one domain to access and use resources located in another domain
- This allows web application running Javascript or HTML5 to access resources in an S3 bucket without using a proxy server
- For AWS, this (commonly) means that a web applications hosted in one s3 bucket can access resources in another S3 bucket.

Storage Transit Services

There are 5 option available here

1) **Single-operational Upload**

- It is the traditional approach for upload where you upload the file in one part
- A single operation upload can upload a file upto 5GB in size, however any file over 100MB should use multipart upload

2) **Multipart Upload**

- It allows you to upload a single object as a set of parts
- Allows for uploading parts of a file concurrently
- Allows for stopping/resuming file uploads
- If transmission of any part fails, you can retransmit that part without affecting other parts.
- After all parts of your objects are uploaded, amazon S3 assembles these parts and creates an object
- Required for objects of 5GB and large, and **Highly suggested** for use when objects are 100MB and larger
- Can be used to upload a file up to 5TB in size

3) **AWS Import/Export**

- It gives the ability to take on-premise data and physically snail mail it to AWS (using a device that you own)
- AWS will import the data either to S3, EBS, or Glacier within one business day of the physical device arriving at AWS

Benefits:

➢ Off-site backup policy
➢ Quickly migrate LARGE amounts of data to the cloud (upto 16TB per job)
➢ Disaster Recovery (AWS will even take s3 data and ship it back to you)

4) **Snowball**

- Snowball is a petabyte-scale data transport solution
- It uses an AWS provided secure transfer appliance
- Quickly move large amounts of data into and out of the AWS cloud

5) **Storage Gateway**

- Connects local data center software appliances to cloud based storage such as Amazon S3.

- Gateway-cached volumes
  → create storage volumes and mount them as iscsi devices on the on-premise servers
  → the gateway will store the data written to this volume in amazon s3 and will cache frequently access data on-premise in the storage device

- Gateway-stored volumes
  → Stores all the data locally (on-premise) in storage volumes
  → Gateway will periodically take snapshots of the data as incremental backups and stores them on S3